

# U1M4.LW.Access and Join Methods

## Part 1


Shkrabatouskaya Vera

[https://github.com/VeraShkrabatouskaya/DataMola\\_Data-Camping-2022](https://github.com/VeraShkrabatouskaya/DataMola_Data-Camping-2022)

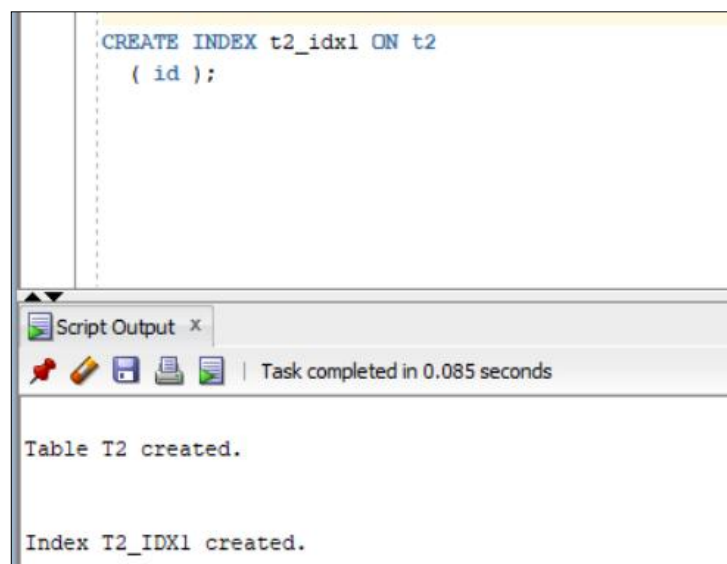
### 1. Table access full scan

#### 1.1. Task 1: Full Scans and the High-water Mark and Block reading

##### Step 1: Create table t2

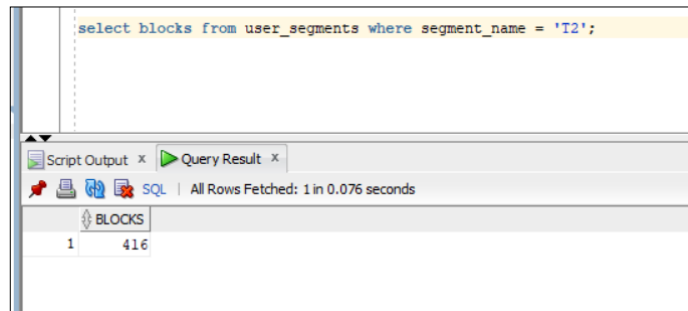


##### Step 2: Create index



### Step 3:

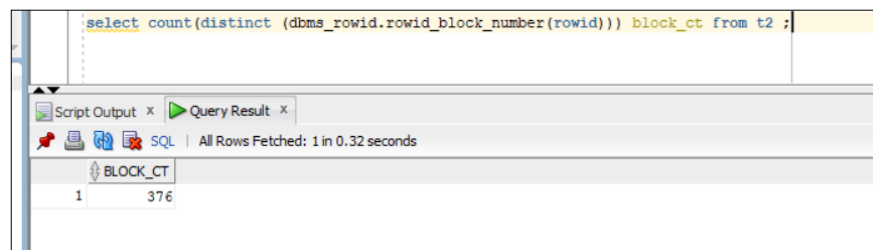
Block count:



The screenshot shows a SQL query in the editor: `select blocks from user_segments where segment_name = 'T2';`. The query result is displayed in a table with one row and two columns: `BLOCKS` and a value of 416. The status bar indicates "All Rows Fetched: 1 in 0.076 seconds".

BLOCKS
416

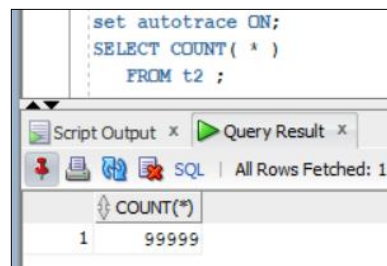
Used Block Count:



The screenshot shows a SQL query in the editor: `select count(distinct (dbms_rowid.rowid_block_number(rowid))) block_ct from t2 ;`. The query result is displayed in a table with one row and two columns: `BLOCK_CT` and a value of 376. The status bar indicates "All Rows Fetched: 1 in 0.32 seconds".

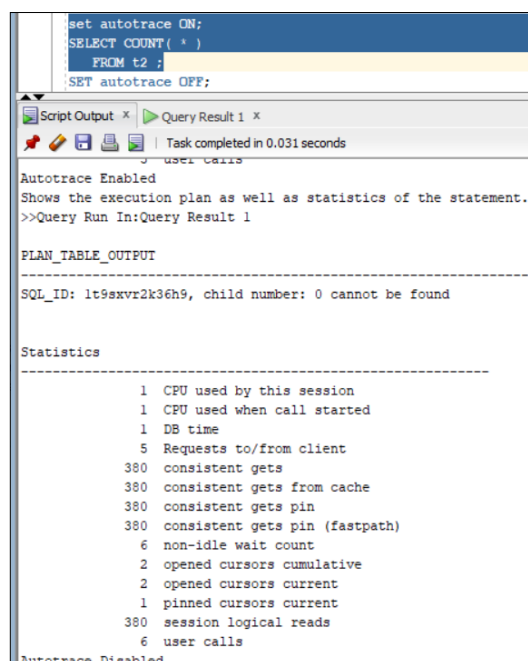
BLOCK_CT
376

Explain Plan / Count rows:



The screenshot shows a SQL query in the editor: `set autotrace ON; SELECT COUNT( * ) FROM t2 ;`. The query result is displayed in a table with one row and two columns: `COUNT(*)` and a value of 99999. The status bar indicates "All Rows Fetched: 1 in 0.031 seconds".

COUNT(*)
99999



The screenshot shows the Autotrace output for the query. It includes the execution plan and statistics. The status bar indicates "Task completed in 0.031 seconds".

Autotrace Enabled  
Shows the execution plan as well as statistics of the statement.  
>>Query Run In:Query Result 1

PLAN\_TABLE\_OUTPUT

SQL\_ID: 1t9sxxvr2k36h9, child number: 0 cannot be found

Statistics

Statistics
1 CPU used by this session
1 CPU used when call started
1 DB time
5 Requests to/from client
380 consistent gets
380 consistent gets from cache
380 consistent gets pin
380 consistent gets pin (fastpath)
6 non-idle wait count
2 opened cursors cumulative
2 opened cursors current
1 pinned cursors current
380 session logical reads
6 user calls

Autotrace Disabled

#### Step 4: Delete All Rows from table

```
DELETE FROM t2;
```

Script Output x Query Result 1 x

Task completed in 0.484 seconds

```
380 consistent gets pin (fastpath)
6 non-idle wait count
2 opened cursors cumulative
2 opened cursors current
1 pinned cursors current
380 session logical reads
6 user calls
Autotrace Disabled
99,999 rows deleted.
```

#### Step 5: Repeat Step 3 and collect results.

Block count:

```
select blocks from user_segments where segment_name = 'T2';

select count(distinct (dbms_rowid.rowid_block_number(rowid))) block_ct from t2 ;
```

Script Output x Query Result x

All Rows Fetched: 1 in 0.005 seconds

BLOCKS
1 416

Used Block Count:

```
select count(distinct (dbms_rowid.rowid_block_number(rowid))) block_ct from t2 ;
```

Script Output x Query Result x

All Rows Fetched: 1 in 0.006 seconds

BLOCK_CT
1 0

Explain Plan / Count rows:

```
SET autotrace ON;
SELECT COUNT( * )
FROM t2 ;
SET autotrace OFF;
```

Script Output x Query Result x

All Rows Fetched: 1 in 0.01 seconds

COUNT(*)
1 0

The screenshot shows the SQL Developer interface with a script named 'VeraDB\_Lab4\_Task\_1.sql'. The script contains the following SQL commands:

```
SET autotrace ON;
SELECT COUNT( * )
  FROM t2 ;
SET autotrace OFF;
```

The 'Script Output' pane shows the execution results:

```
Task completed in 0.043 seconds

Autotrace Enabled
Shows the execution plan as well as statistics of the statement.
>>Query Run In:Query Result

PLAN_TABLE_OUTPUT
-----
SQL_ID: 1t9sxxvr2k36h9, child number: 0 cannot be found

Statistics
-----
      1 CPU used by this session
      1 CPU used when call started
      5 Requests to/from client
     380 consistent gets
     380 consistent gets from cache
     380 consistent gets pin
     380 consistent gets pin (fastpath)
      6 non-idle wait count
      2 opened cursors cumulative
      2 opened cursors current
      1 pinned cursors current
     380 session logical reads
      6 user calls

Autotrace Disabled
```

Step 6: Insert 1 row

The screenshot shows the SQL Developer interface with an SQL script containing the following commands:

```
INSERT INTO t2
 ( ID, T_PAD )
VALUES
 ( 1, '1' );

COMMIT;
```

The 'Script Output' pane shows the execution results:

```
Task completed in 0.038 seconds

1 row inserted.

Commit complete.
```

Step 7: Repeat Step 3 and collect results.

Block count:

The screenshot shows the SQL Developer interface with a script containing the following SQL commands:

```
select blocks from user_segments where segment_name = 'T2';

select count(distinct (dbms_rowid.rowid_block_number(rowid))) block_ct from t2 ;
```

The 'Script Output' pane shows the execution results:

```
Task completed in 0.005 seconds
All Rows Fetched: 1 in 0.005 seconds
```

The 'Query Result' pane displays the following table:

BLOCKS
1

Used Block Count:

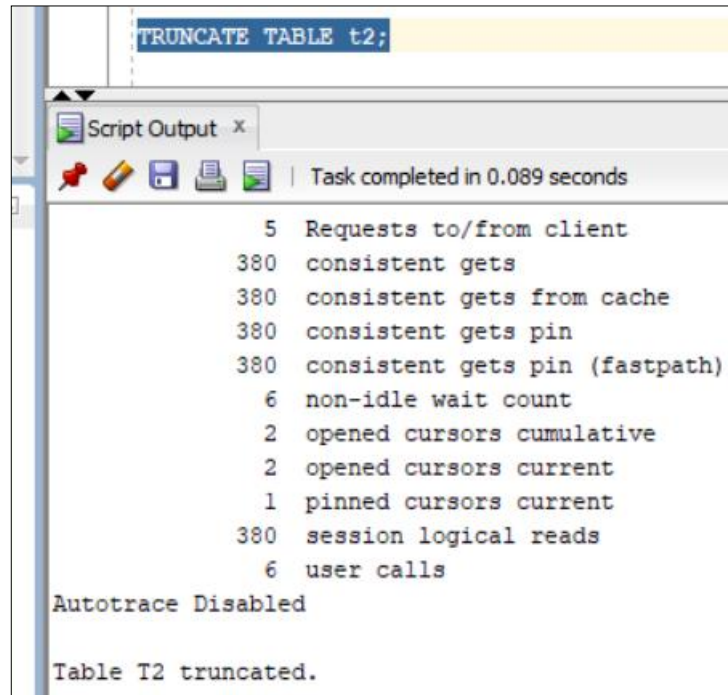
<pre>select count(distinct (dbms_rowid.rowid_block_number(rowid))) block_ct from t2 ;</pre>	
Script Output x	Query Result x
All Rows Fetched: 1 in 0.014 seconds	
BLOCK_CT	
1	1

Explain Plan / Count rows:

<pre>SET autotrace ON; SELECT COUNT( * ) FROM t2 ;</pre>	
Script Output x	Query Result x
All Rows Fetched: 1 in 0.01 seconds	
COUNT(*)	
1	1

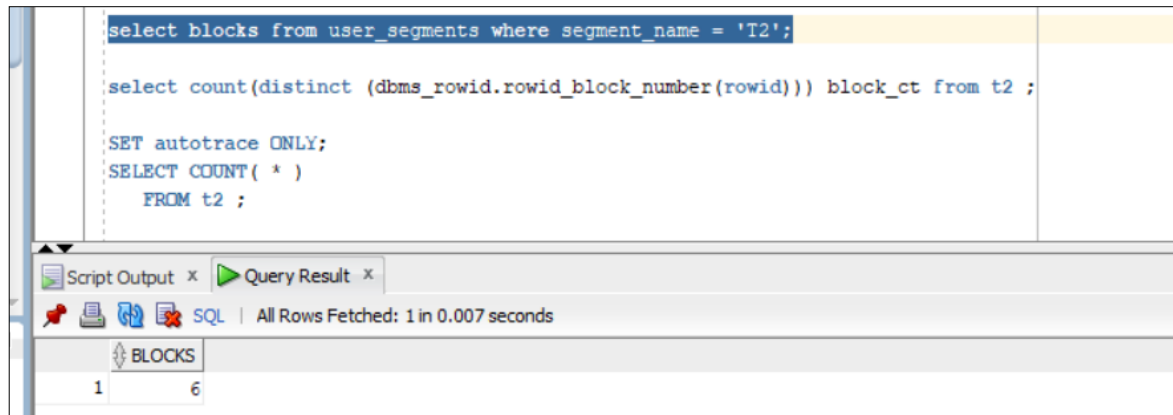
<pre>SET autotrace ON; SELECT COUNT( * ) FROM t2 ; SET autotrace OFF;</pre>	
Script Output x	Query Result x
Task completed in 0.027 seconds	
PLAN_TABLE_OUTPUT	
-----	
SQL_ID: 1t9sxvr2k36h9, child number: 0 cannot be found	
Statistics	
-----	
1	CPU used by this session
1	CPU used when call started
1	DB time
5	Requests to/from client
380	consistent gets
380	consistent gets from cache
380	consistent gets pin
380	consistent gets pin (fastpath)
6	non-idle wait count
2	opened cursors cumulative
2	opened cursors current
1	pinned cursors current
380	session logical reads
6	user calls
Autotrace Disabled	

## Step 8: Truncate Table

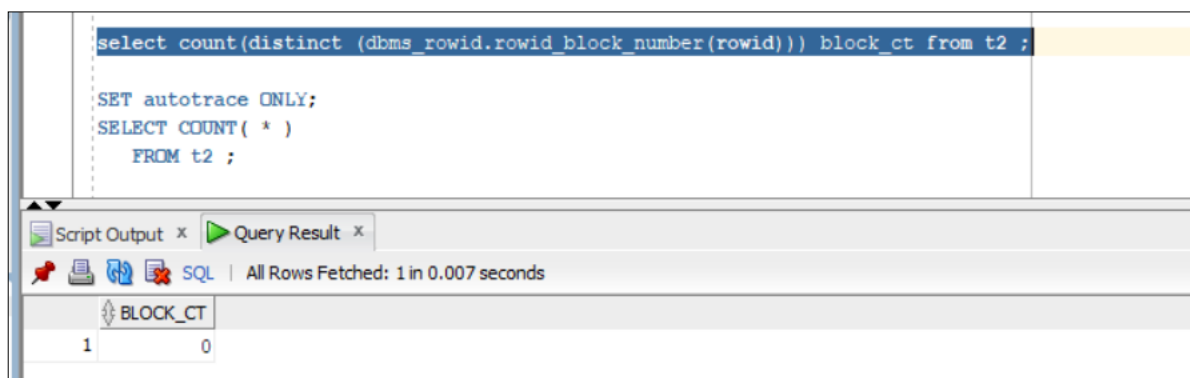


## Step 9: Repeat Step 3 and collect results.

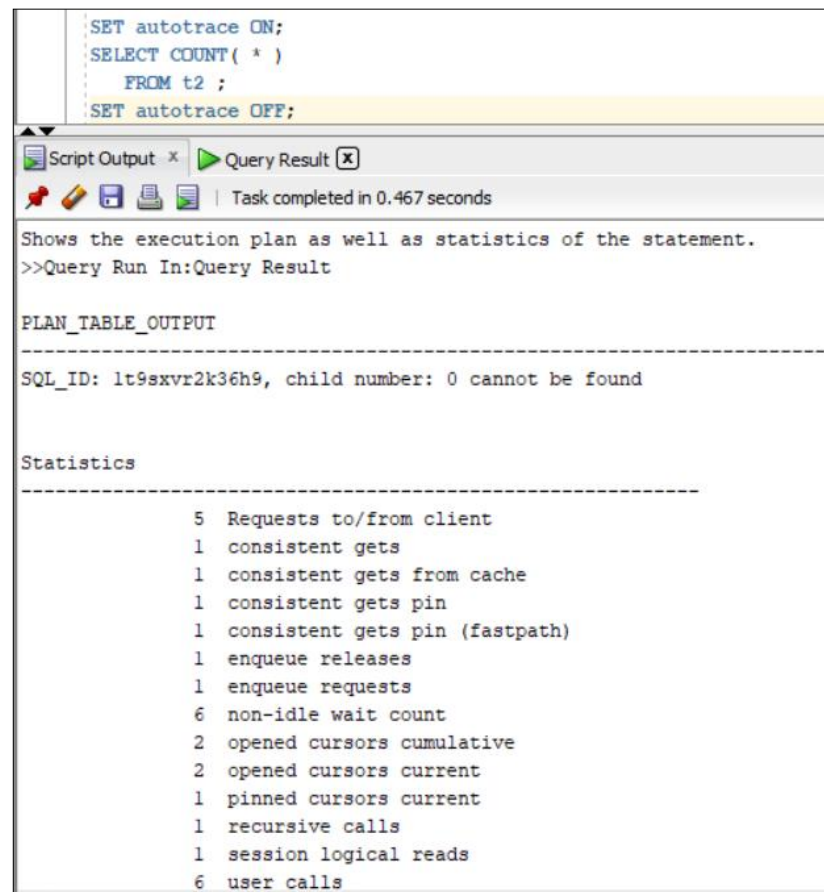
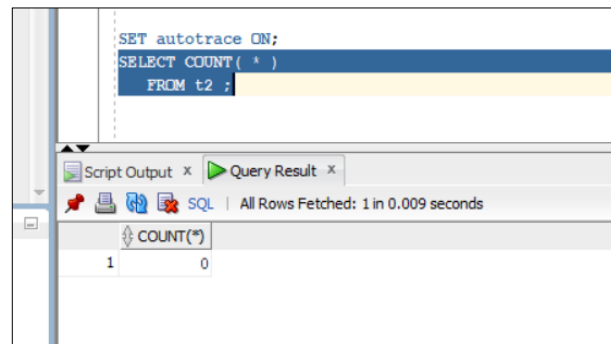
Block count:



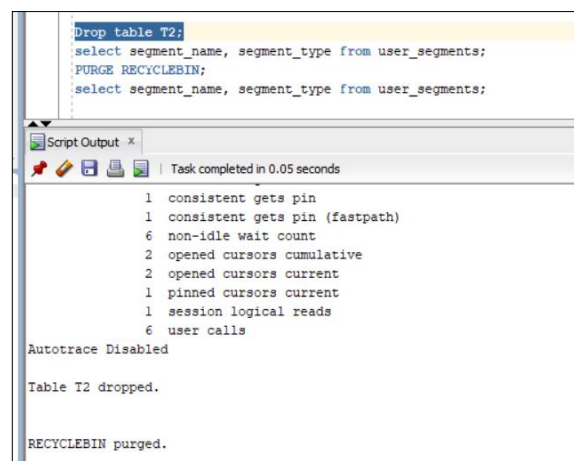
Used Block Count:



Explain Plan / Count rows:



Drop table and purge the recycle bin



Summary table with all result and text description of analyses this results.

Nº	Count of Blocks	Count of Used Blocks	Count of Rows	Consistent gets	Description
1	416	376	99999	380	Full Table
2	416	0	0	380	Empty Table
3	416	1	1	380	Table with 1 row
4	6	0	0	1	Truncate Table

When full scanning an object, all the blocks associated with that object must be retrieved and processed to determine if rows in a block match your query's needs. Oracle must read an entire block into memory in order to get to the row data stored in that block. So, when a full scan occurs, there are actually two things the optimizer needs to consider: how many blocks must be read and how much data in each block will be thrown away.

This example shows that regardless of whether the table is full, empty, or with 1 row, we would need 380 consistent gets from cache for a full table scan, which is inefficient from a cost and performance optimization perspective for tables with 1 row and empty. Even though almost all the rows have been deleted and some blocks have actually become totally unused, the highwater mark remains the same. When a full scan operation occurs, all blocks up to the highwater mark will be read in and scanned, even if they are empty. For tables that are frequently loaded and unloaded (using DELETE instead of TRUNCATE), we may discover that response time suffers.



## Code: Task 1

/\*1.1. Task 1: Full Scans and the High-water Mark and Block reading\*/

CREATE TABLE t2 AS

SELECT TRUNC( rownum / 100 ) id, rpad( rownum,100 ) t\_pad  
FROM dual  
CONNECT BY rownum < 100000;

CREATE INDEX t2\_idx1 ON t2  
( id );

select blocks from user\_segments where segment\_name = 'T2';  
select count(distinct (dbms\_rowid.rowid\_block\_number(rowid))) block\_ct from t2 ;

set autotrace ON;  
SELECT COUNT( \* )  
FROM t2 ;  
SET autotrace OFF;

DELETE FROM t2;

select blocks from user\_segments where segment\_name = 'T2';  
select count(distinct (dbms\_rowid.rowid\_block\_number(rowid))) block\_ct from t2 ;

SET autotrace ON;  
SELECT COUNT( \* )  
FROM t2 ;  
SET autotrace OFF;

INSERT INTO t2  
( ID, T\_PAD )  
VALUES  
( 1,'1' );

COMMIT;

select blocks from user\_segments where segment\_name = 'T2';  
select count(distinct (dbms\_rowid.rowid\_block\_number(rowid))) block\_ct from t2 ;

SET autotrace ON;  
SELECT COUNT( \* )  
FROM t2 ;  
SET autotrace OFF;

TRUNCATE TABLE t2;

select blocks from user\_segments where segment\_name = 'T2';  
select count(distinct (dbms\_rowid.rowid\_block\_number(rowid))) block\_ct from t2 ;

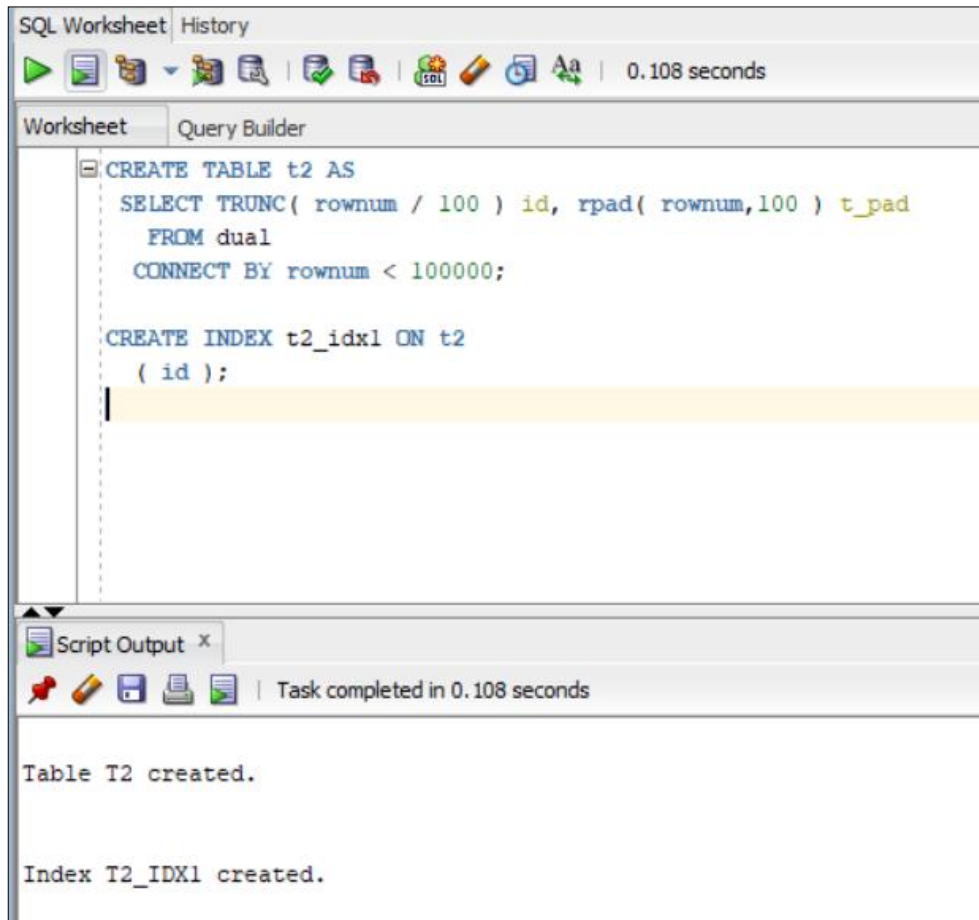
SET autotrace ON;  
SELECT COUNT( \* )  
FROM t2 ;  
SET autotrace OFF;

Drop table T2;  
select segment\_name, segment\_type from user\_segments;  
PURGE RECYCLEBIN;  
select segment\_name, segment\_type from user\_segments;

## 2. Index Scan types

### 2.1. Task 2: Index Clustering factor parameter

Step 1: Create table t2 as on task 1 step 1-2



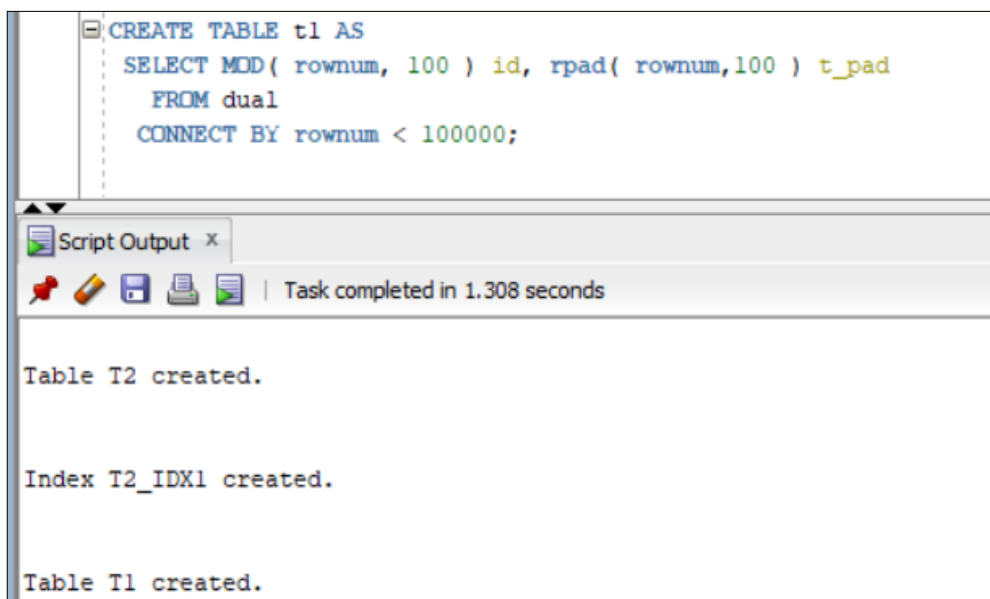
The screenshot shows an SQL Worksheet interface. The top bar includes 'SQL Worksheet' and 'History' tabs, a toolbar with various icons, and a timer showing '0.108 seconds'. Below the toolbar is a 'Worksheet' tab and a 'Query Builder' tab. The main area contains the following SQL code:

```
CREATE TABLE t2 AS
SELECT TRUNC( rownum / 100 ) id, rpad( rownum,100 ) t_pad
FROM dual
CONNECT BY rownum < 100000;

CREATE INDEX t2_idx1 ON t2
( id );
```

Below the code editor is a 'Script Output' window. It shows the execution results: 'Table T2 created.' and 'Index T2\_IDX1 created.' The window also indicates 'Task completed in 0.108 seconds'.

Step 2: Create table t1



The screenshot shows the same SQL Worksheet interface. The main area contains the following SQL code:

```
CREATE TABLE t1 AS
SELECT MOD( rownum, 100 ) id, rpad( rownum,100 ) t_pad
FROM dual
CONNECT BY rownum < 100000;
```

Below the code editor is a 'Script Output' window. It shows the execution results: 'Table T2 created.', 'Index T2\_IDX1 created.', and 'Table T1 created.' The window also indicates 'Task completed in 1.308 seconds'.

### Step 3: Create index

```
CREATE INDEX t1_idx1 ON t1
( id );
```

Script Output x

Task completed in 0.1 seconds

Table T2 created.

Index T2\_IDX1 created.

Table T1 created.

Index T1\_IDX1 created.

### Step 4: Calculate statistic for both tables:

```
EXEC dbms_stats.gather_table_stats( USER,'t1',method_opt=>'FOR ALL COLUMNS SIZE 1',CASCADE=>TRUE );
EXEC dbms_stats.gather_table_stats( USER,'t2',method_opt=>'FOR ALL COLUMNS SIZE 1',CASCADE=>TRUE );
```

Script Output x

Task completed in 0.149 seconds

Table T1 created.

Index T1\_IDX1 created.

PL/SQL procedure successfully completed.

PL/SQL procedure successfully completed.

### Block count:

```
select blocks from user_segments where segment_name = 'T2';
```

Query Result x

All Rows Fetched: 1 in 0.101 seconds

BLOCKS
1 416

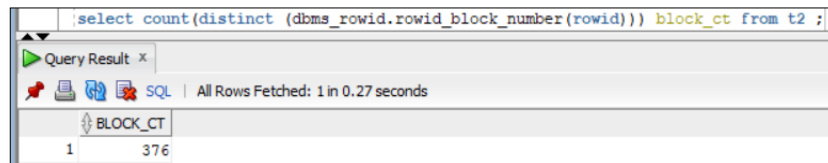
```
select blocks from user_segments where segment_name = 'T1';
```

Query Result x

All Rows Fetched: 1 in 0.053 seconds

BLOCKS
1 384

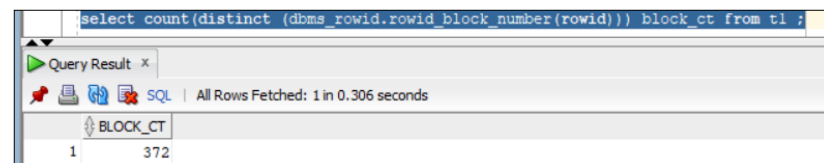
## Used Block Count:



Query Result x

SQL | All Rows Fetched: 1 in 0.27 seconds

BLOCK_CT
1 376

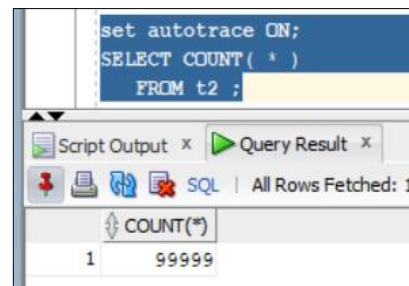


Query Result x

SQL | All Rows Fetched: 1 in 0.306 seconds

BLOCK_CT
1 372

## Explain Plan / Count rows:

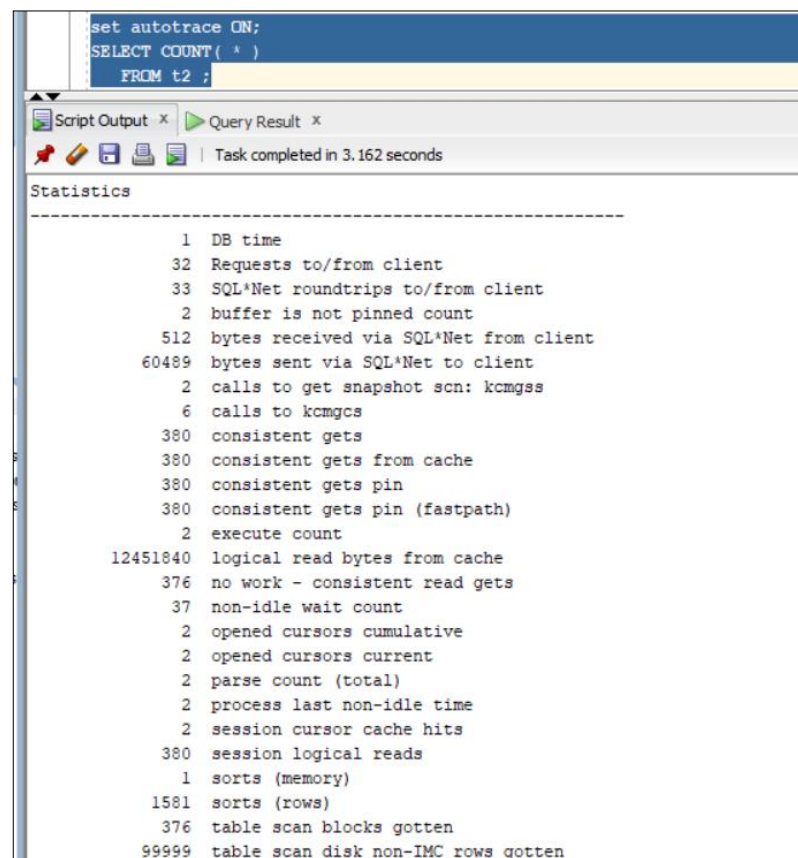


set autotrace ON;  
SELECT COUNT( \* )  
FROM t2 ;

Script Output x Query Result x

SQL | All Rows Fetched: 1

COUNT(*)
1 99999



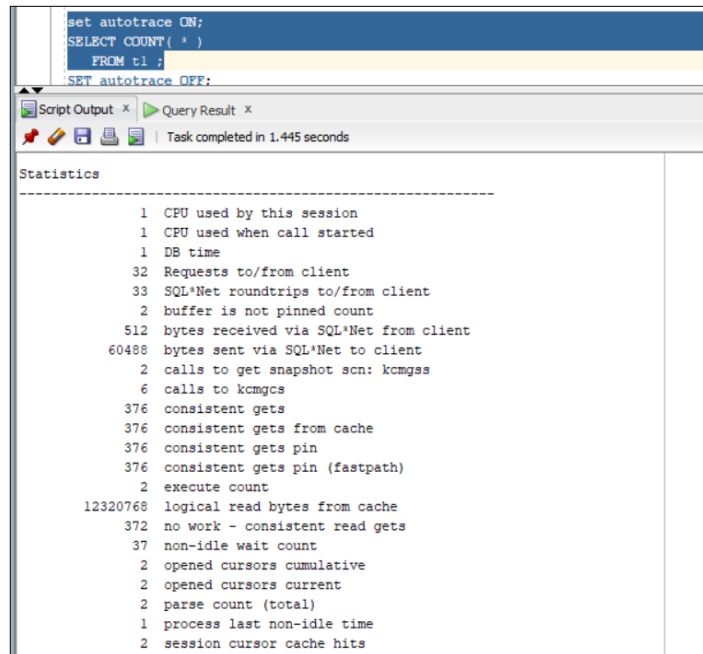
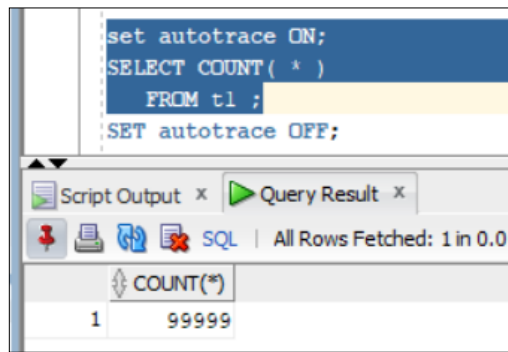
set autotrace ON;  
SELECT COUNT( \* )  
FROM t2 ;

Script Output x Query Result x

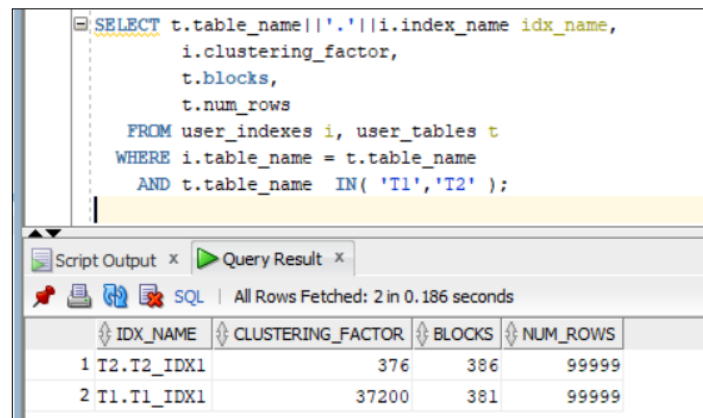
Task completed in 3.162 seconds

Statistics

1	DB time
32	Requests to/from client
33	SQL*Net roundtrips to/from client
2	buffer is not pinned count
512	bytes received via SQL*Net from client
60489	bytes sent via SQL*Net to client
2	calls to get snapshot scn: kcmgss
6	calls to kcmgcs
380	consistent gets
380	consistent gets from cache
380	consistent gets pin
380	consistent gets pin (fastpath)
2	execute count
12451840	logical read bytes from cache
376	no work - consistent read gets
37	non-idle wait count
2	opened cursors cumulative
2	opened cursors current
2	parse count (total)
2	process last non-idle time
2	session cursor cache hits
380	session logical reads
1	sorts (memory)
1581	sorts (rows)
376	table scan blocks gotten
99999	table scan disk non-IMC rows gotten



## Step 5: Select Clustering Factor



## Summary:

- The clustering factor is a measure of how well ordered the table data is as related to the indexed values. It is used to check the cost of a table lookup following an index access (multiplying the clustering factor by index's selectivity gives the cost of the operation). The clustering factor

records the number of blocks that will be read when scanning the index. If the index being used has a large clustering factor, then more table data blocks have to be visited to get the rows in each index block (because adjacent rows are in different blocks). If the clustering factor is close to the number of blocks in the table, then the index is well ordered, but if the clustering factor is close to the number of rows in the table, then the index is not well ordered.

The clustering factor is computed by the following:

- The index is scanned in order.
- The block portion of the ROWID pointed at by the current indexed value is compared to the previous indexed value (comparing adjacent rows in the index).
- If the ROWIDs point to different TABLE blocks, the clustering factor is incremented (this is done for the entire index).

When the clustering factor is higher, the cost of index range scan is higher.

- We have different values for indexes `t1_idx1` and `t2_idx1`. Recall that index entries are stored in sorted order while table data is stored in random order. The clustering factor of an index indicates to the optimizer if data rows containing the same indexed values will be located in the same or a small set of contiguous blocks, or if rows will be scattered across numerous table blocks. We can notice that an index for table T2 would have a lower clustering factor. Lower numbers that are closer to the number of table blocks are used to indicate highly ordered, or clustered, rows of data based on the indexed value. The clustering factor for table T1, however, would be higher and typically closer to the number of rows in the table.
- We can see that index `t2_idx1` has best selective performance in execution `Select clause filtered by IN ( , list of values, )`.

## Code: Task 2

/\*2.1. Task 2: Index Clustering factor parameter\*/

```
CREATE TABLE t2 AS
SELECT TRUNC( rownum / 100 ) id, rpad( rownum,100 ) t_pad
FROM dual
CONNECT BY rownum < 100000;
```

```
CREATE INDEX t2_idx1 ON t2
( id );
```

```
CREATE TABLE t1 AS
SELECT MOD( rownum, 100 ) id, rpad( rownum,100 ) t_pad
FROM dual
CONNECT BY rownum < 100000;
```

```
CREATE INDEX t1_idx1 ON t1
( id );
```

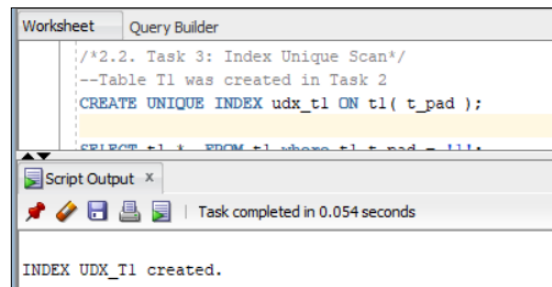
```
EXEC dbms_stats.gather_table_stats( USER,'t1',method_opt=>'FOR ALL COLUMNS SIZE 1',CASCADE=>TRUE );
```

```
EXEC dbms_stats.gather_table_stats( USER,'t2',method_opt=>'FOR ALL COLUMNS SIZE 1',CASCADE=>TRUE );
```


```
SELECT t.table_name || '.' || i.index_name idx_name,
       i.clustering_factor,
       t.blocks,
       t.num_rows
FROM user_indexes i, user_tables t
WHERE i.table_name = t.table_name
AND t.table_name IN( 'T1','T2' );
```

## 2.2. Task 3: Index Unique Scan

### Step 1: Create Unique Index



### Step 2: Data results



Drop table and purge the recycle bin

### Summary:

An index unique scan is chosen when a predicate contains a condition using a column defined with a UNIQUE or PRIMARY KEY index. These types of indexes guarantee that only one row will ever be returned for a specified value. In this cases, the index structure will be traversed from root to leaf block to a single entry, retrieve the rowid, and use it to access the table data block containing the one row.

Column t1.t\_pad has a string value created by the RPAD function, which returns a string value of 100 padded\_length.

Oracle read t1.t\_pad in step 2 as a string with a value of 1 character, but not of 100 characters, so we saw an empty block.

### Code: Task 3

```
/*2.2. Task 3: Index Unique Scan*/  
--Table T1 was created in Task 2  
CREATE UNIQUE INDEX udx_t1 ON t1( t_pad );
```

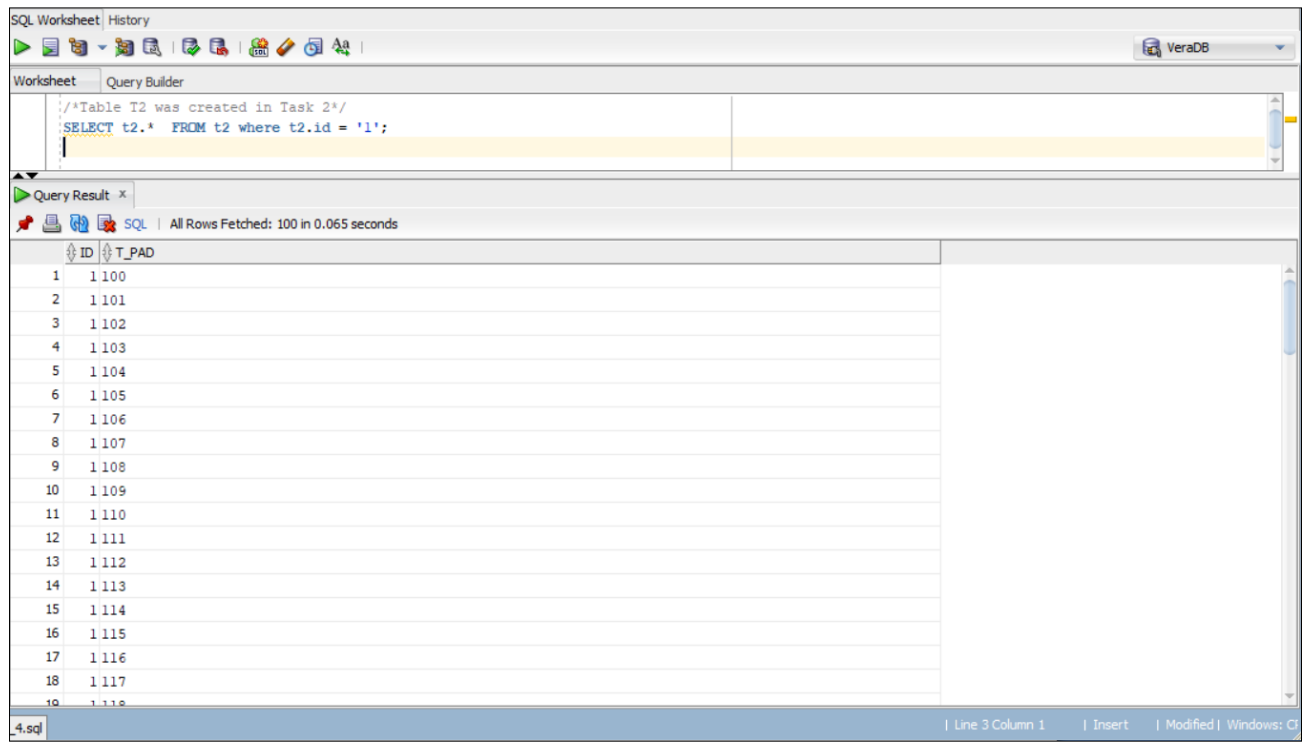
```
SELECT t1.* FROM t1 where t1.t_pad = '1';
```

```
Drop table T1;  
select segment_name, segment_type from user_segments;  
PURGE RECYCLEBIN;  
select segment_name, segment_type from user_segments;
```



## 2.3. Task 4: Index Range Scan

### Step 1: Data results



The screenshot shows a SQL Worksheet interface with a 'Query Builder' tab. The query entered is: `/*Table T2 was created in Task 2*/  
SELECT t2.* FROM t2 where t2.id = '1';` The 'Query Result' tab is active, displaying 19 rows of data. The status bar indicates 'All Rows Fetched: 100 in 0.065 seconds'. The bottom status bar shows 'Line 3 Column 1 | Insert | Modified | Windows: C'.

ID	T_PAD
1	1 100
2	1 101
3	1 102
4	1 103
5	1 104
6	1 105
7	1 106
8	1 107
9	1 108
10	1 109
11	1 110
12	1 111
13	1 112
14	1 113
15	1 114
16	1 115
17	1 116
18	1 117
19	1 118

Drop table and purge the recycle bin

### Summary:

An index range scan is chosen when a predicate contains a condition that will return a range of data.

The index can be unique or non-unique as it is the condition that determines whether or not multiple rows will be returned or not.

Oracle read all t2.t\_pad values (from 100 to 199) for t2.id = 1 in step 1.

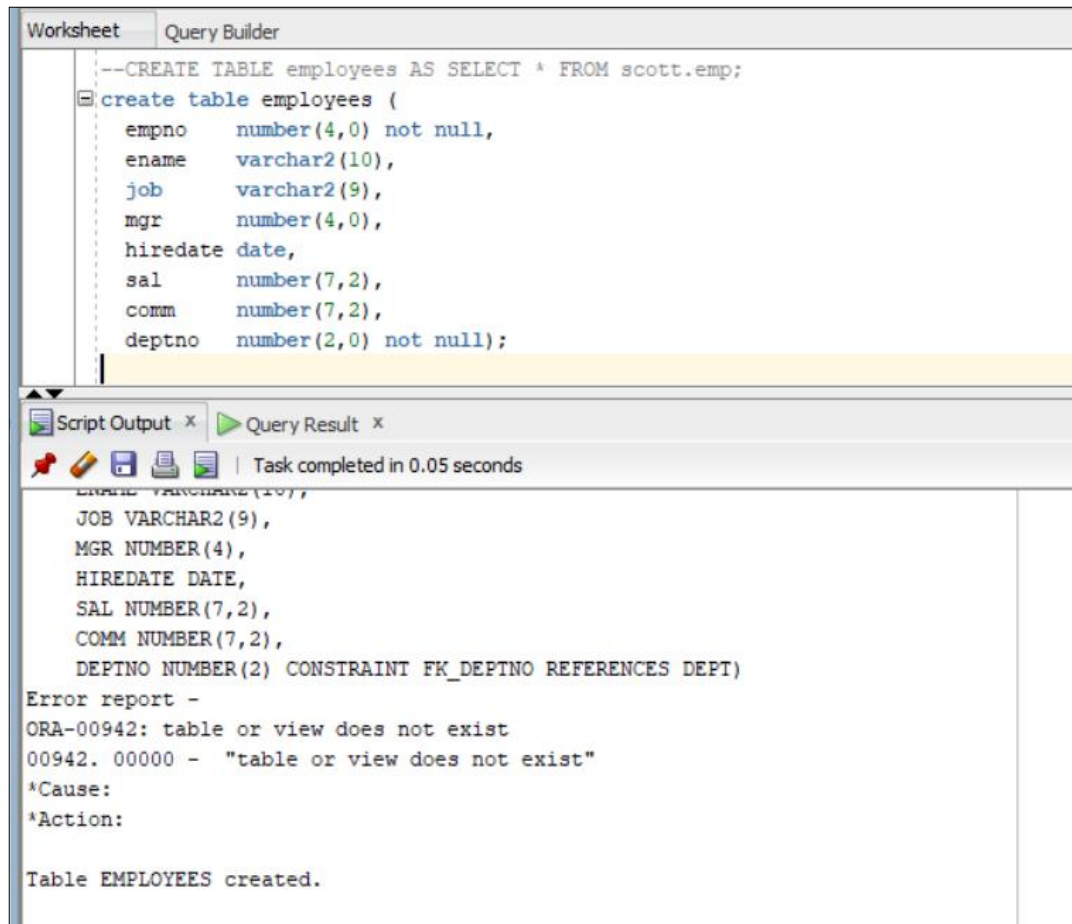
### Code: Task 4

```
/*2.3. Task 4: Index Range Scan*/  
--Table T2 was created in Task 2  
SELECT t2.* FROM t2 where t2.id = '1';
```

```
Drop table T2;  
select segment_name, segment_type from user_segments;  
PURGE RECYCLEBIN;  
select segment_name, segment_type from user_segments;
```

## 2.4. Task 5: Index Skip Scan

### Step 1: Create table



The screenshot shows the SQL Developer interface with the 'Query Builder' tab active. The SQL editor contains the following code:

```
--CREATE TABLE employees AS SELECT * FROM scott.emp;  
create table employees (  
    empno    number(4,0) not null,  
    ename    varchar2(10),  
    job      varchar2(9),  
    mgr      number(4,0),  
    hiredate date,  
    sal      number(7,2),  
    comm     number(7,2),  
    deptno   number(2,0) not null);
```

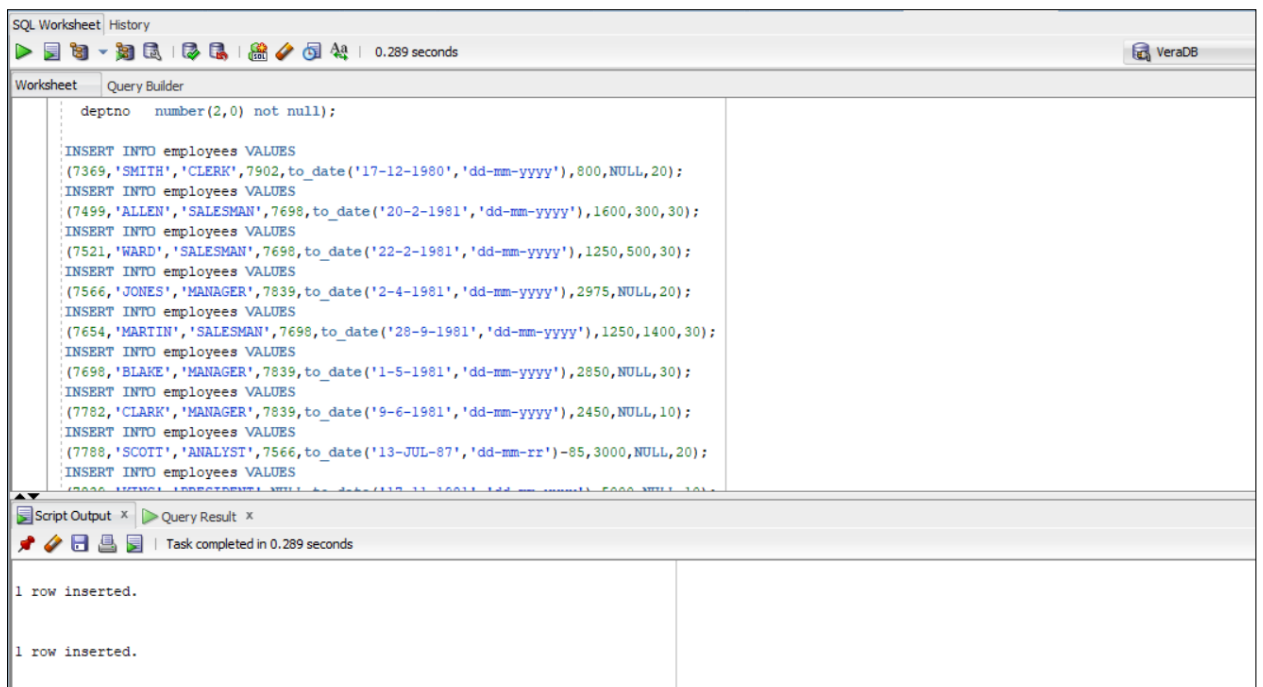
The 'Script Output' tab shows the execution results:

```
TABLE VARCHAR2(10),  
JOB VARCHAR2(9),  
MGR NUMBER(4),  
HIREDATE DATE,  
SAL NUMBER(7,2),  
COMM NUMBER(7,2),  
DEPTNO NUMBER(2) CONSTRAINT FK_DEPTNO REFERENCES DEPT)
```

The 'Query Result' tab shows an error report:

```
Error report -  
ORA-00942: table or view does not exist  
00942. 00000 - "table or view does not exist"  
*Cause:  
*Action:  
  
Table EMPLOYEES created.
```

### Insert values into table



The screenshot shows the SQL Developer interface with the 'Query Builder' tab active. The SQL editor contains the following code:

```
deptno    number(2,0) not null);  
  
INSERT INTO employees VALUES  
(7369,'SMITH','CLERK',7902,to_date('17-12-1980','dd-mm-yyyy'),800,NULL,20);  
INSERT INTO employees VALUES  
(7499,'ALLEN','SALESMAN',7698,to_date('20-2-1981','dd-mm-yyyy'),1600,300,30);  
INSERT INTO employees VALUES  
(7521,'WARD','SALESMAN',7698,to_date('22-2-1981','dd-mm-yyyy'),1250,500,30);  
INSERT INTO employees VALUES  
(7566,'JONES','MANAGER',7839,to_date('2-4-1981','dd-mm-yyyy'),2975,NULL,20);  
INSERT INTO employees VALUES  
(7654,'MARTIN','SALESMAN',7698,to_date('28-9-1981','dd-mm-yyyy'),1250,1400,30);  
INSERT INTO employees VALUES  
(7698,'BLAKE','MANAGER',7839,to_date('1-5-1981','dd-mm-yyyy'),2850,NULL,30);  
INSERT INTO employees VALUES  
(7782,'CLARK','MANAGER',7839,to_date('9-6-1981','dd-mm-yyyy'),2450,NULL,10);  
INSERT INTO employees VALUES  
(7788,'SCOTT','ANALYST',7566,to_date('13-JUL-87','dd-mm-rr'),85,3000,NULL,20);  
INSERT INTO employees VALUES  
(7800,'JONES','MANAGER',7839,to_date('12-1-1981','dd-mm-yyyy'),5000,NULL,30);
```

The 'Script Output' tab shows the execution results:

```
1 row inserted.  
  
1 row inserted.
```

## Step 2: Create Index

```
CREATE INDEX idx_emp01 ON employees
( empno, ename, job );
```

Script Output x Query Result x

Task completed in 0.041 seconds

1 row inserted.

Index IDX\_EMP01 created.

## Step 3: Get trace and statistic of explain plan

Block count:

```
select blocks from user_segments where segment_name = 'EMPLOYEES';
```

Query Result x

All Rows Fetched: 1 in 0.011 seconds

BLOCKS
1 6

Used Block Count:

```
select count(distinct (dbms_rowid.rowid_block_number(rowid))) block_ct from employees ;
```

Query Result x

All Rows Fetched: 1 in 0.013 seconds

BLOCK_CT
1 1

Rows count:

```
SELECT COUNT( * ) FROM employees ;
```

Query Result x

All Rows Fetched: 1 in 0.011 seconds

COUNT(*)
1 14

```
SELECT COUNT(distinct job) ct FROM employees ;
```

Script Output x Query Result x

All Rows Fetched: 1 in 0.014 seconds

CT
1 5

Explain Plan:

SELECT /\*+INDEX\_SS(emp idx\_emp01)\*/ emp.\* FROM employees emp where ename = 'SCOTT';

Script Output x

Query Result x

SQL | All Rows Fetched: 1 in 0.018 seconds

	EMPNO	ENAME	JOB	MGR	HIREDATE	SAL	COMM	DEPTNO
1	7788	SCOTT	ANALYST	7566	19-APR-87	3000	(null)	20

SELECT /\*+FULL\*/ emp.\* FROM employees emp WHERE ename = 'SCOTT';

Script Output x

Query Result x

SQL | All Rows Fetched: 1 in 0.011 seconds

	EMPNO	ENAME	JOB	MGR	HIREDATE	SAL	COMM	DEPTNO
1	7788	SCOTT	ANALYST	7566	19-APR-87	3000	(null)	20

SELECT /\*+INDEX\_SS(emp idx\_emp01)\*/ emp.\* FROM employees emp where ename = 'SCOTT';

SELECT /\*+FULL\*/ emp.\* FROM employees emp WHERE ename = 'SCOTT';

Script Output x

Query Result x

Explain Plan x

SQL | 0.053 seconds

TION	OBJECT_NAME	OPTIONS	CARDINALITY	COST
SELECT STATEMENT				1
TABLE ACCESS	EMPLOYEES	BY INDEX ROWID BATCHED		2
INDEX	IDX_EMP01	SKIP SCAN		1
Access Predicates				
ENAME='SCOTT'				

SELECT /\*+FULL\*/ emp.\* FROM employees emp WHERE ename = 'SCOTT';

Script Output x

Query Result x

Explain Plan x

SQL | 0.014 seconds

TION	OBJECT_NAME	OPTIONS	CARDINALITY	COST
SELECT STATEMENT				1
TABLE ACCESS	EMPLOYEES	FULL		4
Filter Predicates				
ENAME='SCOTT'				4
Other XML				

```
SET autotrace ON;
```

```
SELECT /*+INDEX_SS(emp idx_emp01)*/ emp.* FROM employees emp where ename = 'SCOTT';
```

Script Output x Query Result x Query Result 1 x

Task completed in 1.71 seconds

PLAN\_TABLE\_OUTPUT

---

-----

- dynamic statistics used: dynamic sampling (level=2)
- Warning: basic plan statistics not available. These are only collected when:
  - \* hint 'gather\_plan\_statistics' is used for the statement or
  - \* parameter 'statistics\_level' is set to 'ALL', at session or system level

Statistics

---

```

1  DB time
32  Requests to/from client
33  SQL*Net roundtrips to/from client
2  buffer is not pinned count
566 bytes received via SQL*Net from client
60901 bytes sent via SQL*Net to client
2  calls to get snapshot scn: kcmgss
2  calls to kcmgcs
2  consistent gets
2  consistent gets from cache
2  consistent gets pin
2  consistent gets pin (fastpath)
2  execute count
65536 logical read bytes from cache
2  no work - consistent read gets
34  non-idle wait count

```

```
SELECT /*+FULL*/ emp.* FROM employees emp WHERE ename = 'SCOTT';
```

Script Output x Query Result x Query Result 1 x

Task completed in 1.71 seconds

1 - filter("ENAME"='SCOTT')

Note

---

- dynamic statistics used: dynamic sampling (level=2)
- Warning: basic plan statistics not available. These are only collected when:

PLAN\_TABLE\_OUTPUT

---

- \* hint 'gather\_plan\_statistics' is used for the statement or
- \* parameter 'statistics\_level' is set to 'ALL', at session or system level

Statistics

---

```

1  CPU used by this session
1  CPU used when call started
1  DB time
33  Requests to/from client
33  SQL*Net roundtrips to/from client
2  buffer is not pinned count
547 bytes received via SQL*Net from client
60904 bytes sent via SQL*Net to client
2  calls to get snapshot scn: kcmgss
4  calls to kcmgcs
5  consistent gets
5  consistent gets from cache
5  consistent gets pin

```

Drop table and purge the recycle bin

### Summary:

- An index skip scan is chosen when the predicate contains a condition on a non-leading column in an index and the leading columns are fairly distinct. A skip scan works by logically splitting a multi-column index into smaller subindexes. The number of logical subindexes is determined by the number of distinct values in the leading columns of the index. Therefore, the more distinct the leading columns are, the more logical subindexes would need to be created. If too many subindexes would be required, the operation won't be as efficient as simply doing a full scan. However, in the cases where the number of subindexes needed would be smaller, the operation can be many times more efficient than a full scan as scanning smaller index blocks can be more efficient than scanning larger table blocks.
- Summary table with all result and text description of analyses this results.

Nº	Count of Blocks	Count of Used Blocks	Count of Rows	Consistent gets	Cost	Description
1	6	1	14	5	4	Full Scan: /*+FULL*/
2	6	1	14	2	2	Index Skip Scan: /*+INDEX_SS(emp idx_emp01)*/

In this example, a full table scan is more expensive than an Index Skip Scan, including the fact that a full scan requires 5 consistent gets from cache, versus 2 consistent gets from cache for an Index Skip Scan.

As you can see, an Index Skip Scan is much more efficient. What happened was that the index was logically divided into 5 subindexes (by job) and each subindex was scanned for a match for ename = 'SCOTT'.

For this index scan type, the fewer distinct values the leading column (or columns) have, the fewer logical subindexes will be needed and therefore the fewer total block accesses required.

### Code: Task 5

```
/*2.4. Task 5: Index Skip Scan*/
--CREATE TABLE employees AS SELECT * FROM scott.emp;
create table employees (
  empno  number(4,0) not null,
```

```

ename varchar2(10),
job varchar2(9),
mgr number(4,0),
hiredate date,
sal number(7,2),
comm number(7,2),
deptno number(2,0) not null);

```

```

INSERT INTO employees VALUES
(7369,'SMITH','CLERK',7902,to_date('17-12-1980','dd-mm-yyyy'),800,NULL,20);
INSERT INTO employees VALUES
(7499,'ALLEN','SALESMAN',7698,to_date('20-2-1981','dd-mm-yyyy'),1600,300,30);
INSERT INTO employees VALUES
(7521,'WARD','SALESMAN',7698,to_date('22-2-1981','dd-mm-yyyy'),1250,500,30);
INSERT INTO employees VALUES
(7566,'JONES','MANAGER',7839,to_date('2-4-1981','dd-mm-yyyy'),2975,NULL,20);
INSERT INTO employees VALUES
(7654,'MARTIN','SALESMAN',7698,to_date('28-9-1981','dd-mm-yyyy'),1250,1400,30);
INSERT INTO employees VALUES
(7698,'BLAKE','MANAGER',7839,to_date('1-5-1981','dd-mm-yyyy'),2850,NULL,30);
INSERT INTO employees VALUES
(7782,'CLARK','MANAGER',7839,to_date('9-6-1981','dd-mm-yyyy'),2450,NULL,10);
INSERT INTO employees VALUES
(7788,'SCOTT','ANALYST',7566,to_date('13-JUL-87','dd-mm-rr')-85,3000,NULL,20);
INSERT INTO employees VALUES
(7839,'KING','PRESIDENT',NULL,to_date('17-11-1981','dd-mm-yyyy'),5000,NULL,10);
INSERT INTO employees VALUES
(7844,'TURNER','SALESMAN',7698,to_date('8-9-1981','dd-mm-yyyy'),1500,0,30);
INSERT INTO employees VALUES
(7876,'ADAMS','CLERK',7788,to_date('13-JUL-87','dd-mm-rr')-51,1100,NULL,20);
INSERT INTO employees VALUES
(7900,'JAMES','CLERK',7698,to_date('3-12-1981','dd-mm-yyyy'),950,NULL,30);
INSERT INTO employees VALUES
(7902,'FORD','ANALYST',7566,to_date('3-12-1981','dd-mm-yyyy'),3000,NULL,20);
INSERT INTO employees VALUES
(7934,'MILLER','CLERK',7782,to_date('23-1-1982','dd-mm-yyyy'),1300,NULL,10);

```

```

CREATE INDEX idx_emp01 ON employees
( empno, ename, job );

```

```

select blocks from user_segments where segment_name = 'EMPLOYEES';

```

```

select count(distinct (dbms_rowid.rowid_block_number(rowid))) block_ct from employees ;

```

```

SELECT COUNT( * ) FROM employees ;
SELECT COUNT(distinct job) ct FROM employees ;

```

```

SET autotrace ON;
SELECT /*+INDEX_SS(emp idx_emp01)*/ emp.* FROM employees emp where ename = 'SCOTT';
SET autotrace OFF;

```

```

SET autotrace ON;
SELECT /*+FULL*/ emp.* FROM employees emp WHERE ename = 'SCOTT';
SET autotrace OFF;

```

```

Drop table employees;
select segment_name, segment_type from user_segments;
PURGE RECYCLEBIN;
select segment_name, segment_type from user_segments;

```