

U1M5.LW.Access and Join Methods

Part 2

Shkrabatouskaya Vera

https://github.com/VeraShkrabatouskaya/DataMola_Data-Camping-2022

1. Auto Trace & Explain Plan

1.1. Task 1: Auto Trace configuration training

Step 1: Learning all possible variants of SQL plus utilities autotrace

№	Auto Trace Configuration Options	Description
1	<code>set autotrace off</code>	No AUTOTRACE report is generated. This is the default. Disables all autotrace.
2	<code>set autotrace on</code>	The AUTOTRACE report includes both the optimizer execution path and the SQL statement execution statistics.
3	<code>set autotrace traceonly</code>	Like SET AUTOTRACE ON, but suppresses the printing of the user's query output, if any. If STATISTICS is enabled, query data is still fetched, but not printed.
4	<code>set autotrace on explain</code>	The AUTOTRACE report shows only the optimizer execution path.
5	<code>set autotrace on statistics</code>	The AUTOTRACE report shows only the SQL statement execution statistics.
6	<code>set autotrace on explain statistics</code>	The AUTOTRACE report includes both the optimizer execution path and the SQL statement execution statistics.
7	<code>set autotrace traceonly explain</code>	The AUTOTRACE report includes the optimizer execution path, but suppresses the printing of the user's query output.
8	<code>set autotrace traceonly statistics</code>	The AUTOTRACE report includes the SQL statement execution statistics, but suppresses the printing of the user's query output.
9	<code>set autotrace traceonly explain statistics</code>	The AUTOTRACE report includes both the optimizer execution path and the SQL statement execution statistics, but suppresses the printing of the user's query output.
10	<code>set autotrace off explain</code>	Autotrace function is disabled.

11	set autotrace off statistics	Autotrace function is disabled.
12	set autotrace off explain statistics	Autotrace function is disabled.

Summary: AUTOTRACE is a facility within SQL*Plus to show us the explain plan of the queries we've executed, and the resources they used. The autotrace provides instantaneous feedback including the returned rows, execution plan, and statistics. The user doesn't need to be concerned about trace file locations and formatting since the output is displayed instantly on the screen.

EXPLAIN – generate explain plan once SQL statement is finished.

STATISTICS – generate usage statistics once SQL statement is finished.

TRACEONLY – fetch all data for executed SQL but don't display the data – very useful for testing purposes of huge queries combinations of all above options.

2. Join Methods

2.1. Task 2: Nested Loops Joins

Step 1: Script to create Oracle's "SCOTT" schema (tables EMP, DEPT, BONUS, SALGRADE)

Step 2: Data results

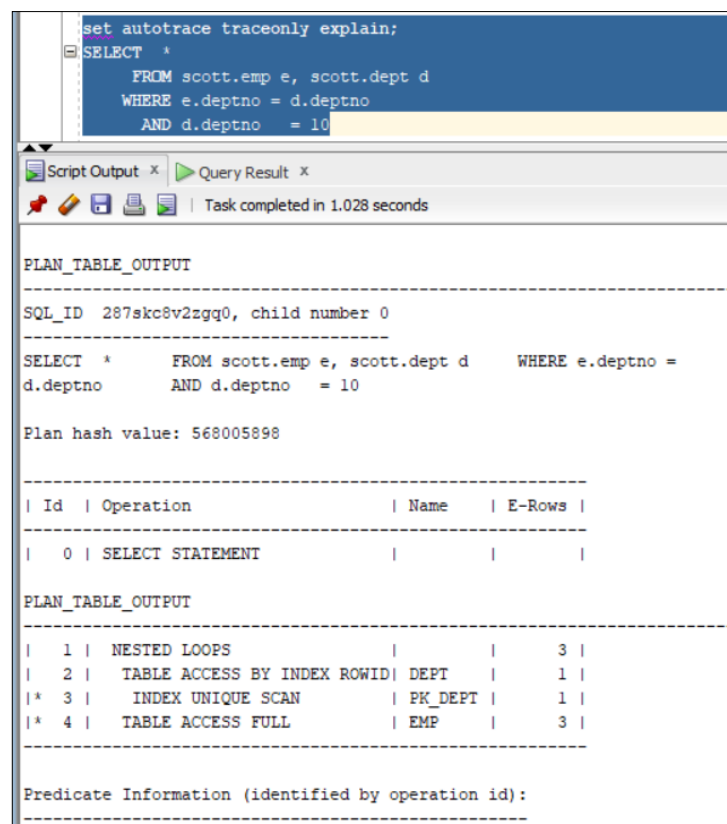
```
SELECT *
FROM emp e, dept d
WHERE e.deptno = d.deptno
AND d.deptno = 10
```

Script Output x Query Result x

SQL | All Rows Fetched: 3 in 4.181 seconds

EMPNO	ENAME	JOB	MGR	HIREDATE	SAL	COMM	DEPTNO	DEPTNO_1	DNAME	LOC
1	7782 CLARK	MANAGER	7839	09-JUN-81	2450	(null)	10	10	ACCOUNTING	NEW YORK
2	7839 KING	PRESIDENT	(null)	17-NOV-81	5000	(null)	10	10	ACCOUNTING	NEW YORK
3	7934 MILLER	CLERK	7782	23-JAN-82	1300	(null)	10	10	ACCOUNTING	NEW YORK

Step 3: Using option Set Autotrace Traceonly Explain



The screenshot shows a SQL Developer window with a query editor containing the following SQL statement:

```
set autotrace traceonly explain;
SELECT *
FROM scott.emp e, scott.dept d
WHERE e.deptno = d.deptno
AND d.deptno = 10
```

Below the editor, the 'Query Result' tab is active, displaying the execution plan for the query. The status bar indicates 'Task completed in 1.028 seconds'.

PLAN_TABLE_OUTPUT

SQL_ID 287skc8v2zqq0, child number 0

SELECT * FROM scott.emp e, scott.dept d WHERE e.deptno = d.deptno AND d.deptno = 10

Plan hash value: 568005898

Id	Operation	Name	E-Rows
0	SELECT STATEMENT		

PLAN_TABLE_OUTPUT

1	NESTED LOOPS		3
2	TABLE ACCESS BY INDEX ROWID	DEPT	1
* 3	INDEX UNIQUE SCAN	PK_DEPT	1
* 4	TABLE ACCESS FULL	EMP	3

Predicate Information (identified by operation id):

```

set autotrace traceonly explain;
SELECT *
FROM scott.emp e, scott.dept d
WHERE e.deptno = d.deptno
AND d.deptno = 10

```

Script Output x Query Result x

Task completed in 1.028 seconds

Statistics

```

-----
1 CPU used by this session
1 CPU used when call started
32 Requests to/from client
33 SQL*Net roundtrips to/from client
5 buffer is not pinned count
585 bytes received via SQL*Net from client
60831 bytes sent via SQL*Net to client
2 calls to get snapshot scn: kcmgss
3 calls to kcmgcs
6 consistent gets
2 consistent gets examination
2 consistent gets examination (fastpath)
6 consistent gets from cache
4 consistent gets pin
4 consistent gets pin (fastpath)
2 execute count
1 index fetch by key
196608 logical read bytes from cache
3 no work - consistent read gets
33 non-idle wait count
2 opened cursors cumulative
2 opened cursors current
2 parse count (total)
1 parse time cpu

```

Set Autotrace Off Explain

In a Nested loops join, we have two tables a driving table and a secondary table. We see that the DEPT table is connected to the EMP table by Nested loops join. The DEPT table is considered first. This connection required 6 consistent gets.

```

explain plan for
SELECT *
FROM scott.emp e, scott.dept d
WHERE e.deptno = d.deptno
AND d.deptno = 10;
select * from table (dbms_xplan.display);

```

Query Result x

All Rows Fetched: 17 in 0.032 seconds

PLAN_TABLE_OUTPUT							
1	Plan hash value: 568005898						
2							
3							
4	Id	Operation	Name	Rows	Bytes	Cost (%CPU)	Time
5							
6	0	SELECT STATEMENT		3	177	5 (0)	00:00:01
7	1	NESTED LOOPS		3	177	5 (0)	00:00:01
8	2	TABLE ACCESS BY INDEX ROWID DEPT		1	20	1 (0)	00:00:01
9	* 3	INDEX UNIQUE SCAN	PK_DEPT	1		0 (0)	00:00:01
10	* 4	TABLE ACCESS FULL	EMP	3	117	4 (0)	00:00:01
11							
12							
13	Predicate Information (identified by operation id):						
14							
15							
16	3	- access("D"."DEPTNO"=10)					
17	4	- filter("E"."DEPTNO"=10)					

If we use the explanation plan for select, we can notice that the cost is 5.

Let's change in the execution plan the type of join method use oracle performance hints. Using the USE_NL hint, let's change the order in which the tables are joined. The table Emp will be considered first.

Using option Set Autotrace Traceonly Explain

The screenshot shows the SQL Developer interface with the 'Query Builder' tab active. The SQL script in the editor is:

```
set autotrace traceonly explain
SELECT /*+ ordered USE_NL(e d) */ *
FROM scott.emp e, scott.dept d
WHERE e.deptno = d.deptno
AND d.deptno = 10
set autotrace off explain;
```

Below the script, the 'Script Output' and 'Query Result' tabs are visible. The 'Query Result' tab shows the execution statistics:

Task completed in 16.256 seconds

Statistics

```
-----
1 CPU used by this session
1 CPU used when call started
32 Requests to/from client
33 SQL*Net roundtrips to/from client
10 buffer is not pinned count
3 buffer is pinned count
611 bytes received via SQL*Net from client
60958 bytes sent via SQL*Net to client
2 calls to get snapshot scn: kcmgcs
3 calls to kcmgcs
9 consistent gets
4 consistent gets examination
4 consistent gets examination (fastpath)
9 consistent gets from cache
5 consistent gets pin
5 consistent gets pin (fastpath)
2 execute count
3 index fetch by key
294912 logical read bytes from cache
4 no work - consistent read gets
33 non-idle wait count
2 opened cursors cumulative
2 opened cursors current
```

Set Autotrace Off Explain

This type of connection requires 9 consecutive gets compared to 6 consecutive gets in the first option.

The screenshot shows the SQL Developer interface with the 'Query Builder' tab active. The SQL script in the editor is:

```
explain plan for
SELECT /*+ ordered USE_NL(e d) */ *
FROM scott.emp e, scott.dept d
WHERE e.deptno = d.deptno
AND d.deptno = 10;
select * from table (dms_xplan.display);
```

Below the script, the 'Script Output' and 'Query Result' tabs are visible. The 'Query Result' tab shows the execution statistics:

All Rows Fetched: 25 in 1.161 seconds

PLAN_TABLE_OUTPUT

Id	Operation	Name	Rows	Bytes	Cost (%CPU)	Time
0	SELECT STATEMENT		3	177	7 (0)	00:00:01
1	NESTED LOOPS		3	177	7 (0)	00:00:01
2	NESTED LOOPS		3	177	7 (0)	00:00:01
3	TABLE ACCESS FULL	EMP	3	117	4 (0)	00:00:01
4	INDEX UNIQUE SCAN	PK_DEPT	1	0	0 (0)	00:00:01
5	TABLE ACCESS BY INDEX ROWID	DEPT	1	20	1 (0)	00:00:01

Predicate Information (identified by operation id):

```
-----
3 - filter("E"."DEPTNO"=10)
4 - access("D"."DEPTNO"=10)
```

Hint Report (identified by operation id / Query Block Name / Object Alias):

```
-----
21 Total hints for statement: 1 (U - Unused (1))
```

The cost of this connection also increases - up to 7 vs. 5 in the first option. Thus, we conclude that Oracle chooses a more favorable option using Nested Loops Joins of tables in terms of cost and performance.

Summary:

Nested loops joins use each row of the query result reached through one access operation to drive into another table. These joins are typically most effective if the result set is limited in size and indexes are present on the columns used for the join. With nested loops, the cost of the operation is based on reading each row of the outer row source and joining it with the matching row of the inner row source.

These kinds of joins are quite robust in that they use very little memory. Since row sets are built one row at a time, there is little overhead required. For that reason, they are actually good for huge result sets except for the fact that building a huge result set one row at a time can take quite a long time.

Code: Task 2

```
/*2.1. Task 2: Nested Loops Joins*/
```

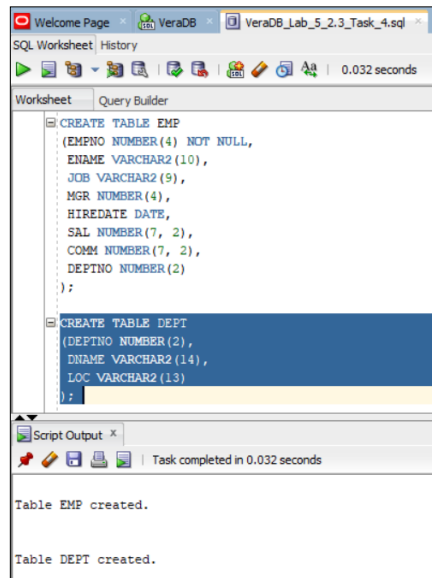
```
SELECT *
  FROM scott.emp e, scott.dept d
 WHERE e.deptno = d.deptno
       AND d.deptno = 10
--
set autotrace traceonly explain;
SELECT *
  FROM scott.emp e, scott.dept d
 WHERE e.deptno = d.deptno
       AND d.deptno = 10
set autotrace off explain;

explain plan for
SELECT *
  FROM scott.emp e, scott.dept d
 WHERE e.deptno = d.deptno
       AND d.deptno = 10;
select * from table (dbms_xplan.display);
--
set autotrace traceonly explain
SELECT /*+ ordered USE_NL(e d) */ *
  FROM scott.emp e, scott.dept d
 WHERE e.deptno = d.deptno
       AND d.deptno = 10
set autotrace off explain;

explain plan for
SELECT /*+ ordered USE_NL(e d) */ *
  FROM scott.emp e, scott.dept d
 WHERE e.deptno = d.deptno
       AND d.deptno = 10;
select * from table (dbms_xplan.display);
--
```

2.2. Task 3: Sort-Merge Joins

Step 1: Create Tables and Insert Values



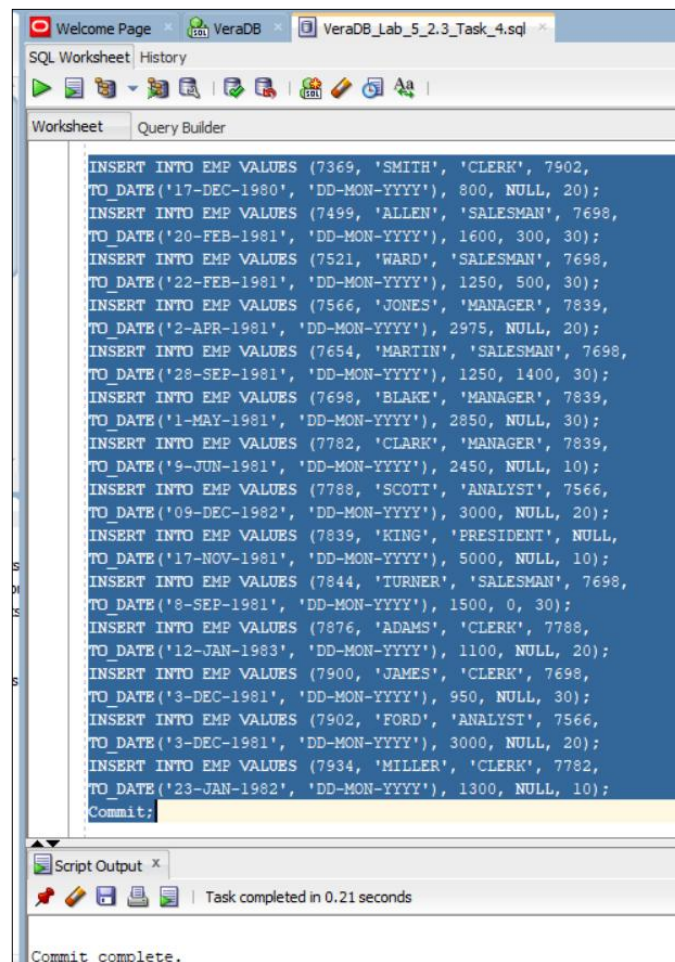
The screenshot shows the SQL Developer interface with the 'Query Builder' tab active. The SQL script in the main window contains two CREATE TABLE statements. The first creates the EMP table with columns EMPNO, ENAME, JOB, MGR, HIREDATE, SAL, COMM, and DEPTNO. The second creates the DEPT table with columns DEPTNO, DNAME, and LOC. The 'Script Output' pane at the bottom shows the execution results: 'Table EMP created.' and 'Table DEPT created.'

```
CREATE TABLE EMP
(EMPNO NUMBER(4) NOT NULL,
ENAME VARCHAR2(10),
JOB VARCHAR2(5),
MGR NUMBER(4),
HIREDATE DATE,
SAL NUMBER(7, 2),
COMM NUMBER(7, 2),
DEPTNO NUMBER(2)
);

CREATE TABLE DEPT
(DEPTNO NUMBER(2),
DNAME VARCHAR2(14),
LOC VARCHAR2(13)
);
```

Table EMP created.

Table DEPT created.



The screenshot shows the SQL Developer interface with the 'Query Builder' tab active. The SQL script in the main window contains a series of INSERT INTO EMP VALUES statements, each followed by a TO_DATE function call. The statements insert data for various employees, including SMITH, ALLEN, WARD, JONES, MARTIN, BLAKE, CLARK, SCOTT, KING, TURNER, ADAMS, JAMES, FORD, and MILLER. The script ends with a COMMIT statement. The 'Script Output' pane at the bottom shows the execution results: 'Commit complete.'

```
INSERT INTO EMP VALUES (7369, 'SMITH', 'CLERK', 7902,
TO_DATE('17-DEC-1980', 'DD-MON-YYYY'), 800, NULL, 20);
INSERT INTO EMP VALUES (7499, 'ALLEN', 'SALESMAN', 7698,
TO_DATE('20-FEB-1981', 'DD-MON-YYYY'), 1600, 300, 30);
INSERT INTO EMP VALUES (7521, 'WARD', 'SALESMAN', 7698,
TO_DATE('22-FEB-1981', 'DD-MON-YYYY'), 1250, 500, 30);
INSERT INTO EMP VALUES (7566, 'JONES', 'MANAGER', 7839,
TO_DATE('2-APR-1981', 'DD-MON-YYYY'), 2975, NULL, 20);
INSERT INTO EMP VALUES (7654, 'MARTIN', 'SALESMAN', 7698,
TO_DATE('28-SEP-1981', 'DD-MON-YYYY'), 1250, 1400, 30);
INSERT INTO EMP VALUES (7698, 'BLAKE', 'MANAGER', 7839,
TO_DATE('1-MAY-1981', 'DD-MON-YYYY'), 2850, NULL, 30);
INSERT INTO EMP VALUES (7782, 'CLARK', 'MANAGER', 7839,
TO_DATE('9-JUN-1981', 'DD-MON-YYYY'), 2450, NULL, 10);
INSERT INTO EMP VALUES (7788, 'SCOTT', 'ANALYST', 7566,
TO_DATE('09-DEC-1982', 'DD-MON-YYYY'), 3000, NULL, 20);
INSERT INTO EMP VALUES (7839, 'KING', 'PRESIDENT', NULL,
TO_DATE('17-NOV-1981', 'DD-MON-YYYY'), 5000, NULL, 10);
INSERT INTO EMP VALUES (7844, 'TURNER', 'SALESMAN', 7698,
TO_DATE('8-SEP-1981', 'DD-MON-YYYY'), 1500, 0, 30);
INSERT INTO EMP VALUES (7876, 'ADAMS', 'CLERK', 7788,
TO_DATE('12-JAN-1983', 'DD-MON-YYYY'), 1100, NULL, 20);
INSERT INTO EMP VALUES (7900, 'JAMES', 'CLERK', 7698,
TO_DATE('3-DEC-1981', 'DD-MON-YYYY'), 950, NULL, 30);
INSERT INTO EMP VALUES (7902, 'FORD', 'ANALYST', 7566,
TO_DATE('3-DEC-1981', 'DD-MON-YYYY'), 3000, NULL, 20);
INSERT INTO EMP VALUES (7934, 'MILLER', 'CLERK', 7782,
TO_DATE('23-JAN-1982', 'DD-MON-YYYY'), 1300, NULL, 10);
Commit;
```

Commit complete.

```

INSERT INTO DEPT VALUES (10, 'ACCOUNTING', 'NEW YORK');
INSERT INTO DEPT VALUES (20, 'RESEARCH', 'DALLAS');
INSERT INTO DEPT VALUES (30, 'SALES', 'CHICAGO');
INSERT INTO DEPT VALUES (40, 'OPERATIONS', 'BOSTON');
Commit;

```

Script Output x

Task completed in 0.075 seconds

1 row inserted.

Commit complete.

Step 2: Data results

Script Output x Explain Plan x Query Result x

All Rows Fetched: 14 in 0.007 seconds

	EMPNO	ENAME	JOB	MGR	HIREDATE	SAL	COMM	DEPTNO	DEPTNO_1	DNAME	LOC
1	7369	SMITH	CLERK	7902	17-DEC-80	800	(null)	20	20	RESEARCH	DALLAS
2	7499	ALLEN	SALESMAN	7698	20-FEB-81	1600	300	30	30	SALES	CHICAGO
3	7521	WARD	SALESMAN	7698	22-FEB-81	1250	500	30	30	SALES	CHICAGO
4	7566	JONES	MANAGER	7839	02-APR-81	2975	(null)	20	20	RESEARCH	DALLAS
5	7654	MARTIN	SALESMAN	7698	28-SEP-81	1250	1400	30	30	SALES	CHICAGO
6	7698	BLAKE	MANAGER	7839	01-MAY-81	2850	(null)	30	30	SALES	CHICAGO
7	7782	CLARK	MANAGER	7839	09-JUN-81	2450	(null)	10	10	ACCOUNTING	NEW YORK
8	7788	SCOTT	ANALYST	7566	09-DEC-82	3000	(null)	20	20	RESEARCH	DALLAS
9	7839	KING	PRESIDENT	(null)	17-NOV-81	5000	(null)	10	10	ACCOUNTING	NEW YORK
10	7844	TURNER	SALESMAN	7698	08-SEP-81	1500	0	30	30	SALES	CHICAGO
11	7876	ADAMS	CLERK	7788	12-JAN-83	1100	(null)	20	20	RESEARCH	DALLAS
12	7900	JAMES	CLERK	7698	03-DEC-81	950	(null)	30	30	SALES	CHICAGO
13	7902	FORD	ANALYST	7566	03-DEC-81	3000	(null)	20	20	RESEARCH	DALLAS
14	7934	MILLER	CLERK	7782	23-JAN-82	1300	(null)	10	10	ACCOUNTING	NEW YORK

Step 3: Using software: SQL plus Auto Trace Utility

Explain Plan:

Script Output x Query Result x

Task completed in 0.543 seconds

SQL_ID c3krmnjs4xl44, child number 0

SELECT * FROM emp e, dept d WHERE e.deptno = d.deptno

Plan hash value: 615168685

Id	Operation	Name	E-Rows	OMem	lMem	Used-Mem
0	SELECT STATEMENT					
* 1	HASH JOIN		14	801K	801K	954K (0)

PLAN_TABLE_OUTPUT

2	TABLE ACCESS FULL	DEPT	4			
3	TABLE ACCESS FULL	EMP	14			

Predicate Information (identified by operation id):

1 - access("E"."DEPTNO"="D"."DEPTNO")


```

set autotrace traceonly explain;
SELECT  *
  FROM emp e, dept d
 WHERE e.deptno = d.deptno;
set autotrace off explain;

```

Script Output x Query Result x

Task completed in 0.532 seconds

Statistics

```

-----
1 CPU used by this session
1 CPU used when call started
1 DB time
32 Requests to/from client
33 SQL*Net roundtrips to/from client
4 buffer is not pinned count
547 bytes received via SQL*Net from client
61690 bytes sent via SQL*Net to client
3 calls to get snapshot scn: kcmgss
6 calls to kcmgcs
10 consistent gets
10 consistent gets from cache
10 consistent gets pin
10 consistent gets pin (fastpath)
1 enqueue releases
1 enqueue requests
3 execute count
327680 logical read bytes from cache
6 no work - consistent read gets
33 non-idle wait count

```

We can see that Oracle joined tables with Hash Joins. Let's change the type of join method use oracle performance hints. (USE_MERGE)

```

set autotrace traceonly explain;
SELECT /*+ ordered USE_MERGE (d e) */ *
  FROM emp e, dept d
 WHERE e.deptno = d.deptno;
set autotrace off explain;

```

Script Output x Query Result x

SQL | All Rows Fetched: 14 in 0.01 seconds

	EMPNO	ENAME	JOB	MGR	HIREDATE	SAL	COMM	DEPTNO	DEPTNO_1	DNAME	LOC
1	7782	CLARK	MANAGER	7839	09-JUN-81	2450	(null)	10	10	ACCOUNTING	NEW YORK
2	7839	KING	PRESIDENT	(null)	17-NOV-81	5000	(null)	10	10	ACCOUNTING	NEW YORK
3	7934	MILLER	CLERK	7782	23-JAN-82	1300	(null)	10	10	ACCOUNTING	NEW YORK
4	7566	JONES	MANAGER	7839	02-APR-81	2975	(null)	20	20	RESEARCH	DALLAS
5	7902	FORD	ANALYST	7566	03-DEC-81	3000	(null)	20	20	RESEARCH	DALLAS
6	7876	ADAMS	CLERK	7788	12-JAN-83	1100	(null)	20	20	RESEARCH	DALLAS
7	7369	SMITH	CLERK	7902	17-DEC-80	800	(null)	20	20	RESEARCH	DALLAS
8	7788	SCOTT	ANALYST	7566	09-DEC-82	3000	(null)	20	20	RESEARCH	DALLAS
9	7521	WARD	SALESMAN	7698	22-FEB-81	1250	500	30	30	SALES	CHICAGO
10	7844	TURNER	SALESMAN	7698	08-SEP-81	1500	0	30	30	SALES	CHICAGO
11	7499	ALLEN	SALESMAN	7698	20-FEB-81	1600	300	30	30	SALES	CHICAGO
12	7900	JAMES	CLERK	7698	03-DEC-81	950	(null)	30	30	SALES	CHICAGO
13	7698	BLAKE	MANAGER	7839	01-MAY-81	2850	(null)	30	30	SALES	CHICAGO
14	7654	MARTIN	SALESMAN	7698	28-SEP-81	1250	1400	30	30	SALES	CHICAGO

```

SELECT /*+ ordered USE_MERGE (d e) */ * FROM emp e, dept d
WHERE e.deptno = d.deptno

Plan hash value: 3406566467

```

Id	Operation	Name	E-Rows	OMem	lMem	Used-Mem
0	SELECT STATEMENT					

PLAN_TABLE_OUTPUT

1	MERGE JOIN		14			
2	SORT JOIN		14	2048	2048	(0)
3	TABLE ACCESS FULL	EMP	14			
4	SORT JOIN		4	2048	2048	(0)
5	TABLE ACCESS FULL	DEPT	4			

Predicate Information (identified by operation id):

```

set autotrace traceonly explain;
SELECT /*+ ordered USE_MERGE (d e) */ *
FROM emp e, dept d
WHERE e.deptno = d.deptno;
set autotrace off explain;

```

Script Output x Query Result x

Task completed in 0.543 seconds

Statistics

1	CPU used by this session
1	CPU used when call started
32	Requests to/from client
33	SQL*Net roundtrips to/from client
4	buffer is not pinned count
578	bytes received via SQL*Net from client
61520	bytes sent via SQL*Net to client
6	calls to get snapshot scn: kcmgss
8	calls to kcmgcs
12	consistent gets
12	consistent gets from cache
12	consistent gets pin
12	consistent gets pin (fastpath)
1	enqueue releases
1	enqueue requests
4	execute count
393216	logical read bytes from cache
6	no work - consistent read gets

The consistent gets parameter is higher when using the Merge Joins for this query, so Oracle chooses a more favorable option using Hash Joins of tables in terms of performance

Summary:

Sort-merge joins read the two tables to be joined independently, sorts the rows from each table (but only those rows that meet the conditions for the table in the WHERE clause) in order by the join key, and then merges the sorted rowsets. The sort operations are the expensive part for this join method. For large row sources that won't fit into memory, the sorts will end

up using temporary disk space to complete. This can be quite memory and time-consuming to complete. But once the rowsets are sorted, the merge happens quickly. To merge, the database alternates down the two lists, compares the top rows, discards rows that are earlier in the sort order than the top of the other list, and only returns matching rows.

Code: Task 3

/*2.2. Task 3: Sort-Merge Joins*/

```
CREATE TABLE EMP
(EMPNO NUMBER(4) NOT NULL,
 ENAME VARCHAR2(10),
 JOB VARCHAR2(9),
 MGR NUMBER(4),
 HIREDATE DATE,
 SAL NUMBER(7, 2),
 COMM NUMBER(7, 2),
 DEPTNO NUMBER(2)
);
```

```
CREATE TABLE DEPT
(DEPTNO NUMBER(2),
 DNAME VARCHAR2(14),
 LOC VARCHAR2(13)
);
```

```
INSERT INTO EMP VALUES (7369, 'SMITH', 'CLERK', 7902,
TO_DATE('17-DEC-1980', 'DD-MON-YYYY'), 800, NULL, 20);
INSERT INTO EMP VALUES (7499, 'ALLEN', 'SALESMAN', 7698,
TO_DATE('20-FEB-1981', 'DD-MON-YYYY'), 1600, 300, 30);
INSERT INTO EMP VALUES (7521, 'WARD', 'SALESMAN', 7698,
TO_DATE('22-FEB-1981', 'DD-MON-YYYY'), 1250, 500, 30);
INSERT INTO EMP VALUES (7566, 'JONES', 'MANAGER', 7839,
TO_DATE('2-APR-1981', 'DD-MON-YYYY'), 2975, NULL, 20);
INSERT INTO EMP VALUES (7654, 'MARTIN', 'SALESMAN', 7698,
TO_DATE('28-SEP-1981', 'DD-MON-YYYY'), 1250, 1400, 30);
INSERT INTO EMP VALUES (7698, 'BLAKE', 'MANAGER', 7839,
TO_DATE('1-MAY-1981', 'DD-MON-YYYY'), 2850, NULL, 30);
INSERT INTO EMP VALUES (7782, 'CLARK', 'MANAGER', 7839,
TO_DATE('9-JUN-1981', 'DD-MON-YYYY'), 2450, NULL, 10);
INSERT INTO EMP VALUES (7788, 'SCOTT', 'ANALYST', 7566,
TO_DATE('09-DEC-1982', 'DD-MON-YYYY'), 3000, NULL, 20);
INSERT INTO EMP VALUES (7839, 'KING', 'PRESIDENT', NULL,
TO_DATE('17-NOV-1981', 'DD-MON-YYYY'), 5000, NULL, 10);
INSERT INTO EMP VALUES (7844, 'TURNER', 'SALESMAN', 7698,
TO_DATE('8-SEP-1981', 'DD-MON-YYYY'), 1500, 0, 30);
INSERT INTO EMP VALUES (7876, 'ADAMS', 'CLERK', 7788,
TO_DATE('12-JAN-1983', 'DD-MON-YYYY'), 1100, NULL, 20);
INSERT INTO EMP VALUES (7900, 'JAMES', 'CLERK', 7698,
TO_DATE('3-DEC-1981', 'DD-MON-YYYY'), 950, NULL, 30);
INSERT INTO EMP VALUES (7902, 'FORD', 'ANALYST', 7566,
TO_DATE('3-DEC-1981', 'DD-MON-YYYY'), 3000, NULL, 20);
INSERT INTO EMP VALUES (7934, 'MILLER', 'CLERK', 7782,
TO_DATE('23-JAN-1982', 'DD-MON-YYYY'), 1300, NULL, 10);
Commit;
```

```
INSERT INTO DEPT VALUES (10, 'ACCOUNTING', 'NEW YORK');
INSERT INTO DEPT VALUES (20, 'RESEARCH', 'DALLAS');
INSERT INTO DEPT VALUES (30, 'SALES', 'CHICAGO');
INSERT INTO DEPT VALUES (40, 'OPERATIONS', 'BOSTON');
Commit;
```

```
SELECT *
  FROM emp e, dept d
 WHERE e.deptno = d.deptno;
```

```
set autotrace traceonly explain;
```

```
SELECT *
  FROM emp e, dept d
 WHERE e.deptno = d.deptno;
```

```
set autotrace off explain;
```

```
set autotrace traceonly explain;
```

```
SELECT /*+ ordered USE_MERGE (d e) */ *
  FROM emp e, dept d
 WHERE e.deptno = d.deptno;
```

```
set autotrace off explain;
```

```
Drop table emp;
```

```
Drop table dept;
```

```
select segment_name, segment_type from user_segments;
```

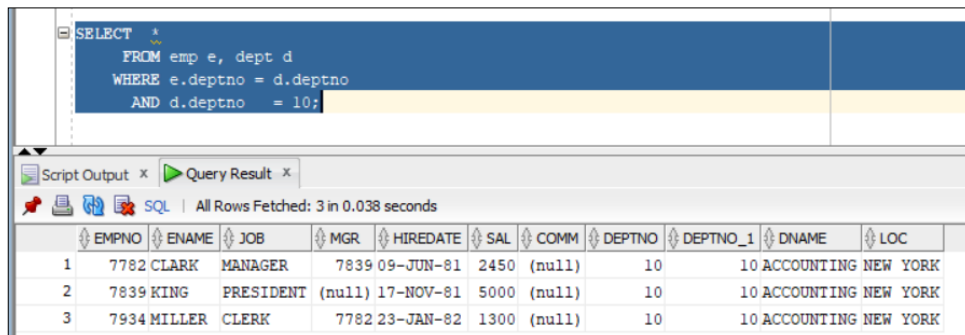
```
PURGE RECYCLEBIN;
```

```
select segment_name, segment_type from user_segments;
```

2.3. Task 4: Hash Joins

Step 1: Create Emp and Dept Tables, Insert Values

Step 2: Data results



The screenshot shows a SQL Developer window with a query editor at the top containing the following SQL statement:

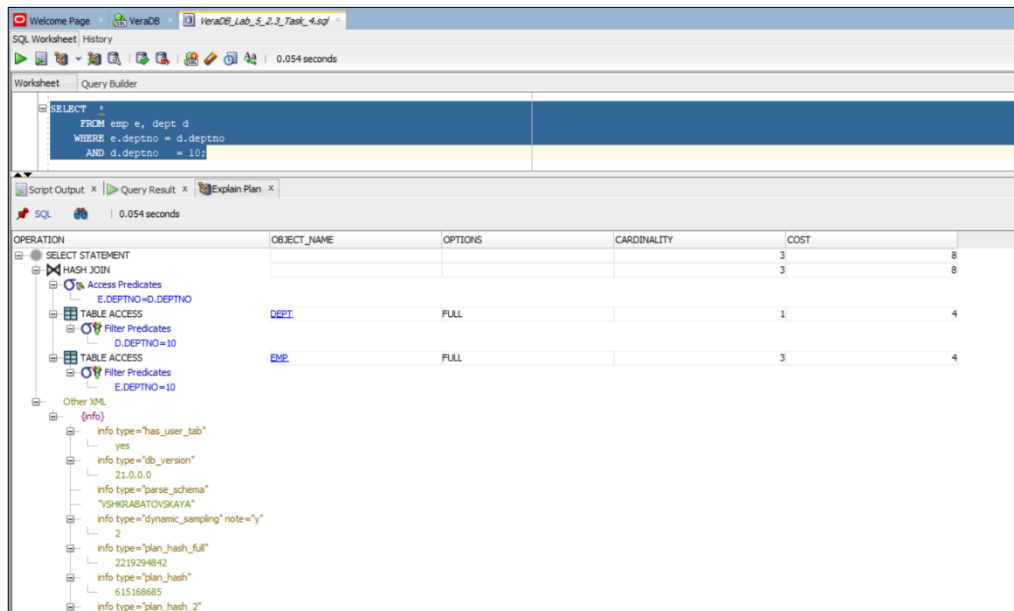
```
SELECT *
FROM emp e, dept d
WHERE e.deptno = d.deptno
AND d.deptno = 10;
```

Below the query editor, the 'Query Result' tab is active, displaying the results of the query. The status bar indicates 'All Rows Fetched: 3 in 0.038 seconds'. The results are shown in a table with 11 columns: EMPNO, ENAME, JOB, MGR, HIREDATE, SAL, COMM, DEPTNO, DEPTNO_1, DNAME, and LOC. There are 3 rows of data.

EMPNO	ENAME	JOB	MGR	HIREDATE	SAL	COMM	DEPTNO	DEPTNO_1	DNAME	LOC
1	7782 CLARK	MANAGER	7839	09-JUN-81	2450	(null)	10	10	ACCOUNTING	NEW YORK
2	7839 KING	PRESIDENT	(null)	17-NOV-81	5000	(null)	10	10	ACCOUNTING	NEW YORK
3	7934 MILLER	CLERK	7782	23-JAN-82	1300	(null)	10	10	ACCOUNTING	NEW YORK

Step 3: Using software: Oracle SQL Developer

Explain Plan:



The screenshot shows the 'Explain Plan' tab in SQL Developer. The query is the same as in Step 2. The explain plan shows a 'HASH JOIN' operation. The plan details are as follows:

OPERATION	OBJECT_NAME	OPTIONS	CARDINALITY	COST
SELECT STATEMENT			3	8
HASH JOIN			3	8
Access Predicates				
E.DEPTNO=D.DEPTNO				
TABLE ACCESS	DEPT	FULL	1	4
Filter Predicates				
D.DEPTNO=10				
TABLE ACCESS	EMP	FULL	3	4
Filter Predicates				
E.DEPTNO=10				
Other XML				
(info)				
info type="has_user_tab"				
yes				
info type="db_version"				
21.0.0.0				
info type="parse_schema"				
"VSHKRABATOVSKAYA"				
info type="dynamic_sampling" note="y"				
2				
info type="plan_hash_full"				
2219294842				
info type="plan_hash"				
615158685				
info type="plan_hash_2"				

We see that in the hash join plan, the smaller hash table DEPT is listed first and the EMP table is listed second. The decision as to which table is smallest depends not just on the number of rows but the size of those rows as well, since the entire row must be stored in the hash table.

If we change the order in which the tables are joined, we see that the cost does not change.

<pre> SELECT /*+ ordered USE_HASH(e d) */ * FROM emp e, dept d WHERE e.deptno = d.deptno AND d.deptno = 10; </pre>				
Explain Plan x SQL 0.024 seconds				
OPERATION	OBJECT_NAME	OPTIONS	CARDINALITY	COST
SELECT STATEMENT				8
HASH JOIN			3	8
Access Predicates				
E.DEPTNO=D.DEPTNO				
TABLE ACCESS	EMP	FULL	3	4
Filter Predicates				
E.DEPTNO=10				
TABLE ACCESS	DEPT	FULL	1	4
Filter Predicates				
D.DEPTNO=10				
Other XML				
(info)				
info type="has_user_tab"				
yes				
info type="db_version"				
21.0.0.0				
info type="parse_schema"				
"VSHRABATOVSKEYA"				
info type="dynamic_sampling" note="Y"				
2				
info type="plan_hash_full"				
2533197573				
info type="plan_hash"				
1123238657				
info type="plan_hash"				

Summary:

HASH joins are the usual choice of the Oracle optimizer when the memory is set up to accommodate them. In a HASH join, Oracle accesses one table (usually the smaller of the joined results) and builds a hash table on the join key in memory. It then scans the other table in the join (usually the larger one) and probes the hash table for matches to it. Hash joins are only possible if the join is an equi-join.

Code: Task 4

/*2.3. Task 4: Hash Joins*/

```

SELECT *
  FROM emp e, dept d
 WHERE e.deptno = d.deptno
        AND d.deptno = 10;

```

```

SELECT /*+ ordered USE_HASH(e d) */ *
  FROM emp e, dept d
 WHERE e.deptno = d.deptno
        AND d.deptno = 10;

```

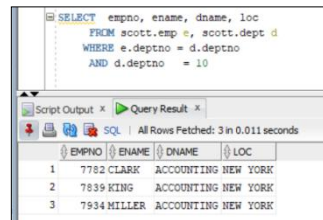
```

Drop table emp;
Drop table dept;
select segment_name, segment_type from user_segments;
PURGE RECYCLEBIN;
select segment_name, segment_type from user_segments;

```

2.4. Task 5: Cartesian Joins

Step 1: Data results

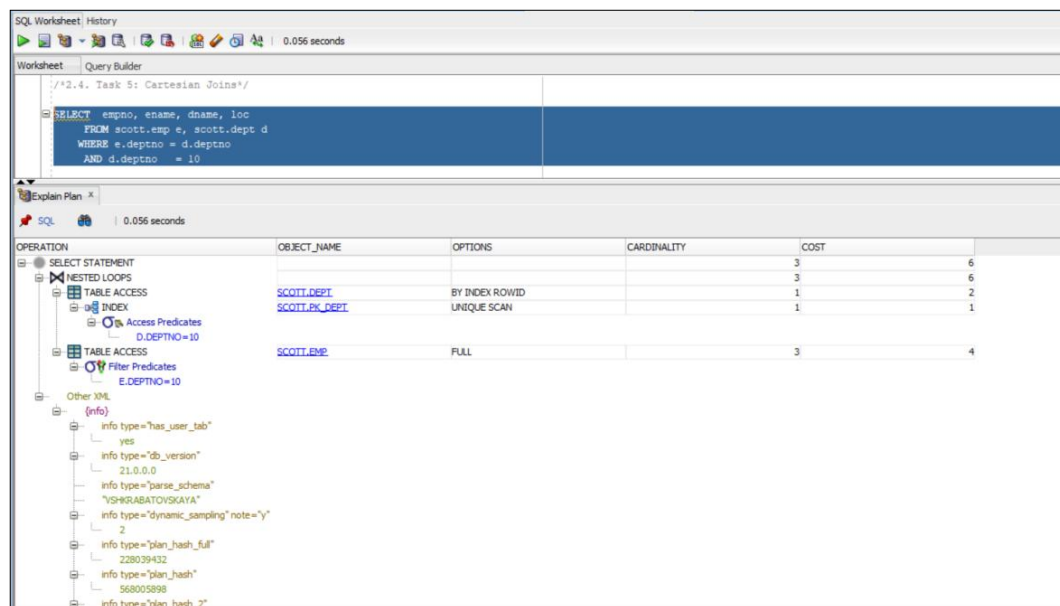


The screenshot shows a SQL query in the 'Script Output' window and its results in the 'Query Result' window. The query is a Cartesian join between SCOTT.EMP and SCOTT.DEPT where both departments are 10.

EMPNO	ENAME	DNAME	LOC
1	7782 CLARK	ACCOUNTING	NEW YORK
2	7839 KING	ACCOUNTING	NEW YORK
3	7934 MILLER	ACCOUNTING	NEW YORK

Step 2: Using software: Oracle SQL Developer

Explain Plan:

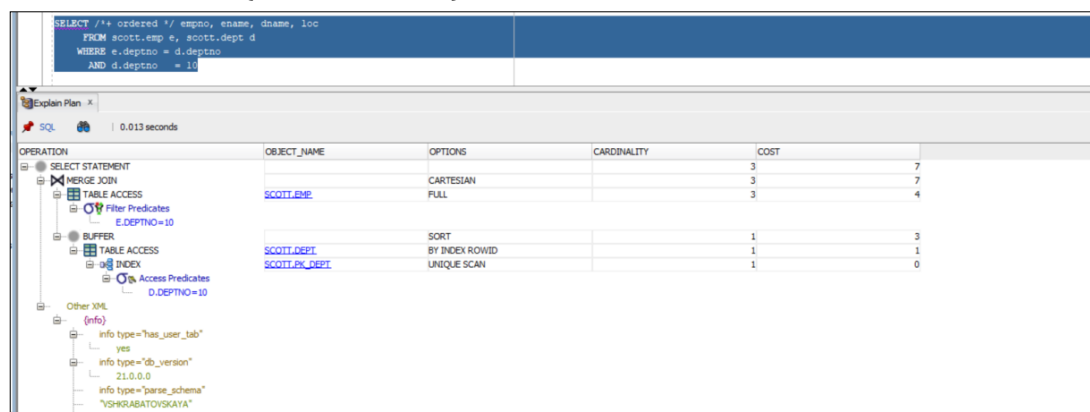


The screenshot shows the 'Explain Plan' window for the query. The execution plan is a NESTED LOOPS join. The first table accessed is SCOTT.DEPT, and the second is SCOTT.EMP. The cost is 6.

OPERATION	OBJECT_NAME	OPTIONS	CARDINALITY	COST
SELECT STATEMENT				6
NESTED LOOPS				6
TABLE ACCESS	SCOTT.DEPT	BY INDEX ROWID	1	2
INDEX	SCOTT.PK_DEPT	UNIQUE SCAN	1	1
TABLE ACCESS	SCOTT.EMP	FULL	3	4

In this query, Oracle suggests using Nested loops joins.

Let's change in the execution plan the type of join method use oracle performance hints. (USE_MERGE)



The screenshot shows the 'Explain Plan' window for the query with the USE_MERGE hint. The execution plan is a MERGE JOIN. The first table accessed is SCOTT.EMP, and the second is SCOTT.DEPT. The cost is 7.

OPERATION	OBJECT_NAME	OPTIONS	CARDINALITY	COST
SELECT STATEMENT				7
MERGE JOIN		CARTESIAN		7
TABLE ACCESS	SCOTT.EMP	FULL	3	4
TABLE ACCESS	SCOTT.DEPT	BY INDEX ROWID	1	3
INDEX	SCOTT.PK_DEPT	UNIQUE SCAN	1	0

The cost parameter is higher when using the Merge Cartesian Joins for this query.

Summary:

Cartesian joins occur when all the rows from one table are joined to all the rows of another table.

Therefore, the total number of rows resulting from the join equals the number of rows from one table (A) multiplied by the number of rows in the other table (B) such that $A \times B = \text{total rows in the result}$

set. Cartesian joins often occur when a join condition is overlooked or left out such that there isn't a specified join column so the only operation possible is to simply join everything from one row source to everything from the other.

Code: Task 5

/*2.4. Task 5: Cartesian Joins*/

```
SELECT empno, ename, dname, loc  
  FROM scott.emp e, scott.dept d  
 WHERE e.deptno = d.deptno  
        AND d.deptno = 10
```

```
SELECT /*+ ordered */ empno, ename, dname, loc  
  FROM scott.emp e, scott.dept d  
 WHERE e.deptno = d.deptno  
        AND d.deptno = 10
```


Step 5: Using software: Oracle SQL Developer

The screenshot shows the Oracle SQL Developer interface. At the top, a SQL window contains the following query:

```
SELECT *
FROM scott.emp e
RIGHT OUTER JOIN scott.dept d
ON e.deptno = d.deptno;
```

Below the SQL window, the 'Explain Plan' tab is active, showing the execution plan for the query. The plan is as follows:

OPERATION	OBJECT_NAME	OPTIONS	CARDINALITY	COST
SELECT STATEMENT			13	8
HASH JOIN		OUTER	13	8
Access Predicates E.DEPTNO(+) = D.DEPTNO				
TABLE ACCESS	SCOTT.DEPT	FULL	4	4
TABLE ACCESS	SCOTT.EMP	FULL	12	4

Below the execution plan, there is an 'Other XML' section with the following information:

```
{info}
  info type="has_user_tab"
  yes
  info type="db_version"
  21.0.0.0
  info type="parse_schema"
  "VSHKRABATOVSKAYA"
  info type="dynamic_sampling" note="y"
```

The cost parameter doesn't change when using the Left Joins or the Right Joins for this query.

Code: Task 6

/*2.5. Task 6: Left/Right Outer Joins*/

--Script to create Oracle's "SCOTT" schema (tables EMP, DEPT, BONUS, SALGRADE)

```
SELECT *
FROM scott.emp e
LEFT OUTER JOIN scott.dept d
ON e.deptno = d.deptno;
```

```
SELECT *
FROM scott.emp e
RIGHT OUTER JOIN scott.dept d
ON e.deptno = d.deptno;
```

2.6. Task 7: Full Outer Join

Step 1: Script to create Oracle's "SCOTT" schema (tables EMP, DEPT, BONUS, SALGRADE)

Step 2: Data results

<pre>SELECT * FROM scott.emp e FULL OUTER JOIN scott.dept d ON e.deptno = d.deptno;</pre>												
Query Result x												
All Rows Fetched: 13 in 0.01 seconds												
	EMPNO	ENAME	JOB	MGR	HIREDATE	SAL	COMM	DEPTNO	DEPTNO_1	DNAME	LOC	
1	7369	SMITH	CLERK	7902	17-DEC-80	800	(null)	20	20	RESEARCH	DALLAS	
2	7499	ALLEN	SALESMAN	7698	20-FEB-81	1600	300	30	30	SALES	CHICAGO	
3	7521	WARD	SALESMAN	7698	22-FEB-81	1250	500	30	30	SALES	CHICAGO	
4	7566	JONES	MANAGER	7839	02-APR-81	2975	(null)	20	20	RESEARCH	DALLAS	
5	7654	MARTIN	SALESMAN	7698	28-SEP-81	1250	1400	30	30	SALES	CHICAGO	
6	7698	BLAKE	MANAGER	7839	01-MAY-81	2850	(null)	30	30	SALES	CHICAGO	
7	7782	CLARK	MANAGER	7839	09-JUN-81	2450	(null)	10	10	ACCOUNTING	NEW YORK	
8	7839	KING	PRESIDENT	(null)	17-NOV-81	5000	(null)	10	10	ACCOUNTING	NEW YORK	
9	7844	TURNER	SALESMAN	7698	08-SEP-81	1500	0	30	30	SALES	CHICAGO	
10	7900	JAMES	CLERK	7698	03-DEC-81	950	(null)	30	30	SALES	CHICAGO	
11	7902	FORD	ANALYST	7566	03-DEC-81	3000	(null)	20	20	RESEARCH	DALLAS	
12	7934	MILLER	CLERK	7782	23-JAN-82	1300	(null)	10	10	ACCOUNTING	NEW YORK	
13	(null)	(null)	(null)	(null)	(null)	(null)	(null)	(null)	40	OPERATIONS	BOSTON	

Step 3: Using software: Oracle SQL Developer

<pre>SELECT * FROM scott.emp e FULL OUTER JOIN scott.dept d ON e.deptno = d.deptno;</pre>												
Query Result x Explain Plan x												
0.014 seconds												
OPERATION	OBJECT_NAME	OPTIONS	CARDINALITY	COST								
SELECT STATEMENT				13	8							
VIEW	SYS_VW_FOJ_0			13	8							
HASH JOIN		FULL OUTER		13	8							
Access Predicates												
E.DEPTNO=D.DEPTNO												
TABLE ACCESS	SCOTT.DEPT	FULL		4	4							
TABLE ACCESS	SCOTT.EMP	FULL		12	4							
Other XML												
(info)												
info type="has_user_tab"												
yes												
info type="db_version"												
21.0.0.0												
info type="parse_schema"												
"VSHKRABATOVSKAYA"												
info type="dynamic_sampling" note="y"												
2												

Code: Task 7

/*2.6. Task 7: Full Outer Join*/

--Script to create Oracle's "SCOTT" schema (tables EMP, DEPT, BONUS, SALGRADE)

```
SELECT *
FROM scott.emp e
FULL OUTER JOIN scott.dept d
ON e.deptno = d.deptno;
```

2.7. Task 8: Semi Joins

Step 1: Create Emp and Dept Tables, Insert Values

Step 2: Data results

```
select /* using in */ ename, job
from emp e
where deptno IN (select deptno from dept d);
```

	ENAME	JOB
1	CLARK	MANAGER
2	KING	PRESIDENT
3	MILLER	CLERK
4	SMITH	CLERK
5	JONES	MANAGER
6	SCOTT	ANALYST
7	ADAMS	CLERK
8	FORD	ANALYST
9	ALLEN	SALESMAN
10	WARD	SALESMAN
11	MARTIN	SALESMAN
12	BLAKE	MANAGER
13	TURNER	SALESMAN
14	JAMES	CLERK

Step 3: Using software: SQL plus Auto Trace Utility

Worksheet

Query Builder

```
set autotrace traceonly explain;
select /* using in */ ename, job
from emp e
where deptno IN (select deptno from dept d);
set autotrace off explain;
```

Script Output x

Query Result x

Task completed in 0.446 seconds

SQL_ID 7y9vandpw4sr, child number 0

select /* using in */ ename, job from emp e where deptno IN (select
deptno from dept d)

Plan hash value: 230627304

Id	Operation	Name	E-Rows	OMem	lMem	Used-Mem
0	SELECT STATEMENT					

PLAN_TABLE_OUTPUT

1	HASH JOIN SEMI		14	863K	863K	982K (0)
2	TABLE ACCESS FULL	EMP	14			
3	TABLE ACCESS FULL	DEPT	4			

Predicate Information (identified by operation id):

1 - access("DEPTNO"="DEPTNO")

Note

Summary: A semi-join is a join between two sets of data (tables) where rows from the first set are returned, based on the presence or absence of at least one matching row in the other set.

The requirements for Oracle's cost based optimizer to decide to use a semi-join:

- The statement must use either the keyword IN (= ANY) or the keyword EXISTS
- The statement must have a subquery in the IN or EXISTS clause
- If the statement uses the EXISTS syntax, it must use a correlated subquery (to get the expected results)
- The IN or EXISTS clause may not be contained inside an OR branch

Code: Task 8

/*2.7. Task 8: Semi Joins*/

```
select /* using in */ ename, job  
from emp e  
where deptno IN (select deptno from dept d);
```

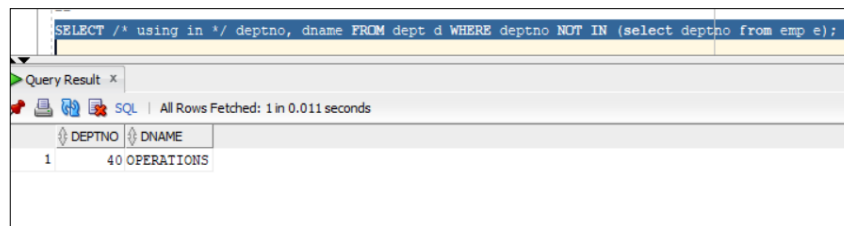
set autotrace traceonly explain;

```
select /* using in */ ename, job  
from emp e  
where deptno IN (select deptno from dept d);  
set autotrace off explain;
```

2.8. Task 9: Anti Joins

Step 1: Create Emp and Dept Tables, Insert Values

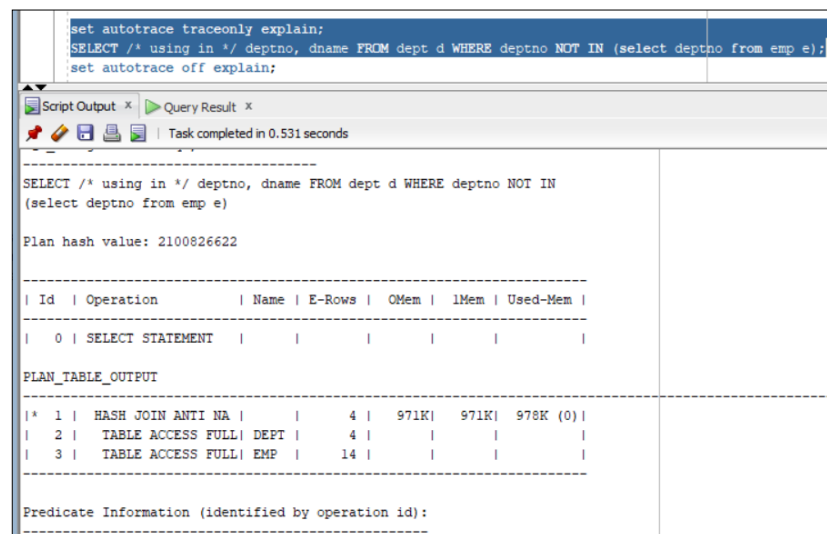
Step 2: Data results



The screenshot shows a SQL Developer window with a query result. The query is: `SELECT /* using in */ deptno, dname FROM dept d WHERE deptno NOT IN (select deptno from emp e);`. The result is displayed in a table with two columns: DEPTNO and DNAME. There is one row with the value 1 in the DEPTNO column and 40 OPERATIONS in the DNAME column.

DEPTNO	DNAME
1	40 OPERATIONS

Step 3: Using software: SQL plus Auto Trace Utility



The screenshot shows the SQL plus Auto Trace Utility output. It includes the query: `SELECT /* using in */ deptno, dname FROM dept d WHERE deptno NOT IN (select deptno from emp e);`. The output shows the plan hash value: 2100826622. Below this is a table with columns: Id, Operation, Name, E-Rows, OMem, lMem, and Used-Mem. The table shows the execution plan for the query, including the SELECT STATEMENT and the TABLE ACCESS FULL operations for DEPT and EMP. The output also includes the PLAN_TABLE_OUTPUT and the Predicate Information (identified by operation id):

Id	Operation	Name	E-Rows	OMem	lMem	Used-Mem
0	SELECT STATEMENT					

Id	Operation	Name	E-Rows	OMem	lMem	Used-Mem
1	HASH JOIN ANTI NA		4	971K	971K	978K (0)
2	TABLE ACCESS FULL	DEPT	4			
3	TABLE ACCESS FULL	EMP	14			

Summary: Anti-joins are basically the same as semi-joins in that they are an optimization option that can be applied to nested loop, hash, and merge joins. However, they are the opposite of semi-joins in terms of the data they return.

Code: Task 9

/*2.8. Task 9: Anti Joins*/

SELECT /* using in */ deptno, dname FROM dept d WHERE deptno NOT IN (select deptno from emp e);

set autotrace traceonly explain;

SELECT /* using in */ deptno, dname FROM dept d WHERE deptno NOT IN (select deptno from emp e);

set autotrace off explain;

2.9. Task 10: Prepare summary table

Compare all possible variant of join methods and join access methods and fill the table below:

Join Access "A"	Join Access "B"	Nested Loop	Hash Join	Sort-Merge Join
Small Table	Small Table	+	+(if the join is an equi-join)	+
Small Table	Indexed Small Table	+	+(if the join is an equi-join)	+
Indexed Small Table	Indexed Small Table	+	+(if the join is an equi-join)	+
Big Table	Big Table	+/- (quite a long time)	+(if the table can fit in memory)	+/- (for large row sources can be quite memory and time-consuming to complete)
Big Table	Small Table	+/- (quite a long time)	+/- (if the smaller table can fit in memory)	+/- (for large row sources can be quite memory and time-consuming to complete)
Big Table	Indexed Small Table	+/- (quite a long time)	+/- (if the smaller table can fit in memory)	+/- (for large row sources can be quite memory and time-consuming to complete)
Big Table	Indexed Big Table	+/- (quite a long time)	+(if the table can fit in memory)	+/- (for large row sources can be quite memory and time-consuming to complete)
Indexed Big Table	Small Table	+/- (quite a long time)	+/- (if the smaller table can fit in memory)	+/- (for large row sources can be quite memory and time-consuming to complete)
Indexed Big Table	Indexed Small Table	+/- (quite a long time)	+/- (if the smaller table can fit in memory)	+
Indexed Big Table	Indexed Big Table	+/- (quite a long time)	+(if the table can fit in memory)	+